

Managing Mixed Modeled Data using Polystore.

Subhasis Dasgupta, Ph.D

A decorative light blue triangle is located in the bottom right corner of the slide.

Mixed Modeled Data

- We produce, use, analyze, and store a massive amount of data for various purposes.
- We modeled them differently to enable analysis.
- Hence we need to handle different model of data in reality.


```
{
  "empId": "SJ011M5",
  "personal": {
    "name": "Smith Jones",
    "gender": "Male",
    "age": 28,
    "address": {
      "streetAddress": "77 24th Street",
      "city": "New York",
      "state": "NY",
      "postalCode": "10038"
    }
  },
  "profile": {
    "designation": "Deputy General",
    "department": "Finance"
  }
}
```




Column or Category

Name	Customer	Age	Month	Year	Lat	Long
John	John	31	April	1980	1294567	1294567
John	John	31	March	1979	7003027	1700302
Peter	Peter	4	August	1988	481708	481708
Sam	Sam	8	May	1980	482790	188279
George	George	20	August	1968	4388746	4888746

Table or Relation



automated data mining survey
responses consumer transcripts
qualitative root cause
classification insights
ad-hoc analysis product
reviews search volume of the
customer dashboards consumer
trends ad-hoc analysis early warning



Model Specific Databases

www.dbdb.io

Database of Databases

Browse

Leaderboards

Recent

Search

Accounts

Refine by

Start Year

Enable

End Year

Enable

Country

☐ Australia

☐ Austria

☐ Bangladesh

Show more

Compatible With

☐ Access

☐ Accumulo

☐ BoltDB

Show more

Embeds / Uses

☐ BadgerDB

☐ Berkeley DB

☐ BoltDB

Show more

Derived From

☐ Accumulo

☐ Adabas

☐ Adaptive Server Enterprise

Show more

Inspired By

☐ ArangoDB

☐ Berkeley DB

☐ BigQuery

Show more

Operating System

☐ AIX

☐ All OS with Java VM

☐ Android

Show more

Programming Languages

☐ ActionScript

☐ Assembly

☐ Bash

Show more

Project Types

☐ Academic


☐ Commercial


Show more


Begin searching!


Search


/ 1 3 4 A B C D E F G H I J K L M N O P Q R S T U V W X Y Z All



iRdb
Last updated Feb. 3, 2021, 11:29 a.m.



1010data
Last updated June 3, 2018, 10:35 p.m.



3store
Last updated July 18, 2019, 6:01 p.m.



4D
Last updated Dec. 10, 2019, 11:54 a.m.



GAMMA
Last updated June 8, 2018, 8:31 a.m.



GaussDB
Last updated Sept. 17, 2020, 8:34 p.m.



GBase
Last updated April 20, 2019, 11:02 a.m.


GonFire
Last updated May 1, 2019, 11:02 a.m.


Pika
Last updated May 8, 2020, 9:26 a.m.


piladb
Last updated June 8, 2018, 11:11 p.m.


Pincaster
Last updated June 5, 2019, 8:15 p.m.

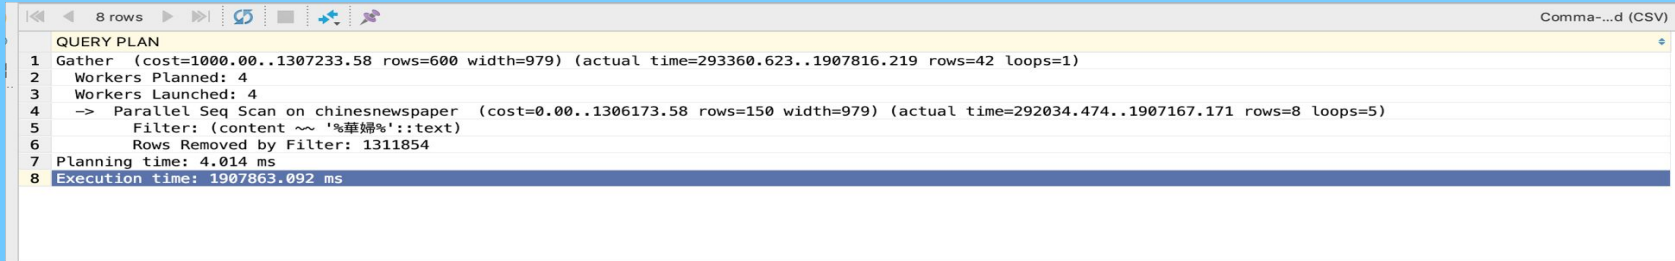

Pinecone
Last updated June 5, 2019, 8:15 p.m.

“No One Size Fits All”

LARGE VOLUME AND CAPABILITIES

EXAMPLE OF LARGE QUERIES

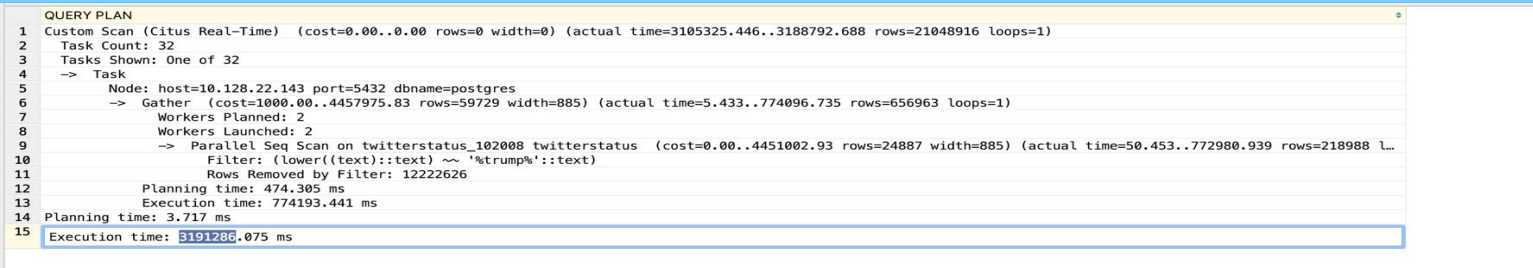
EXPLAIN ANALYZE SELECT * from chinesenewspaper WHERE content LIKE '%華婦%'



The screenshot shows a database interface with a query execution plan for the query: **EXPLAIN ANALYZE SELECT * from chinesenewspaper WHERE content LIKE '%華婦%'**. The plan is displayed in a table with 8 rows. The first row is the header 'QUERY PLAN'. The subsequent rows show the execution details: 1. Gather (cost=1000.00..1307233.58 rows=600 width=979) (actual time=293360.623..1907816.219 rows=42 loops=1); 2. Workers Planned: 4; 3. Workers Launched: 4; 4. -> Parallel Seq Scan on chinesenewspaper (cost=0.00..1306173.58 rows=150 width=979) (actual time=292034.474..1907167.171 rows=8 loops=5); 5. Filter: (content ~ '%華婦% '::text); 6. Rows Removed by Filter: 1311854; 7. Planning time: 4.014 ms; 8. Execution time: 1907863.092 ms.

	QUERY PLAN
1	Gather (cost=1000.00..1307233.58 rows=600 width=979) (actual time=293360.623..1907816.219 rows=42 loops=1)
2	Workers Planned: 4
3	Workers Launched: 4
4	-> Parallel Seq Scan on chinesenewspaper (cost=0.00..1306173.58 rows=150 width=979) (actual time=292034.474..1907167.171 rows=8 loops=5)
5	Filter: (content ~ '%華婦% '::text)
6	Rows Removed by Filter: 1311854
7	Planning time: 4.014 ms
8	Execution time: 1907863.092 ms

EXPLAIN ANALYZE SELECT * FROM twitterstatus where lower(text) Like '%trump%';



The screenshot shows a database interface with a query execution plan for the query: **EXPLAIN ANALYZE SELECT * FROM twitterstatus where lower(text) Like '%trump%';**. The plan is displayed in a table with 15 rows. The first row is the header 'QUERY PLAN'. The subsequent rows show the execution details: 1. Custom Scan (Citrus Real-Time) (cost=0.00..0.00 rows=0 width=0) (actual time=3105325.446..3188792.688 rows=21048916 loops=1); 2. Task Count: 32; 3. Tasks Shown: One of 32; 4. -> Task; 5. Node: host=10.128.22.143 port=5432 dbname=postgres; 6. -> Gather (cost=1000.00..4457975.83 rows=59729 width=885) (actual time=5.433..774096.735 rows=656963 loops=1); 7. Workers Planned: 2; 8. Workers Launched: 2; 9. -> Parallel Seq Scan on twitterstatus_102008 twitterstatus (cost=0.00..4451002.93 rows=24887 width=885) (actual time=50.453..772980.939 rows=218988 loops=1); 10. Filter: (lower((text)::text) ~ '%trump% '::text); 11. Rows Removed by Filter: 12222626; 12. Planning time: 474.305 ms; 13. Execution time: 774193.441 ms; 14. Planning time: 3.717 ms; 15. Execution time: 3191286.075 ms.

	QUERY PLAN
1	Custom Scan (Citrus Real-Time) (cost=0.00..0.00 rows=0 width=0) (actual time=3105325.446..3188792.688 rows=21048916 loops=1)
2	Task Count: 32
3	Tasks Shown: One of 32
4	-> Task
5	Node: host=10.128.22.143 port=5432 dbname=postgres
6	-> Gather (cost=1000.00..4457975.83 rows=59729 width=885) (actual time=5.433..774096.735 rows=656963 loops=1)
7	Workers Planned: 2
8	Workers Launched: 2
9	-> Parallel Seq Scan on twitterstatus_102008 twitterstatus (cost=0.00..4451002.93 rows=24887 width=885) (actual time=50.453..772980.939 rows=218988 loops=1)
10	Filter: (lower((text)::text) ~ '%trump% '::text)
11	Rows Removed by Filter: 12222626
12	Planning time: 474.305 ms
13	Execution time: 774193.441 ms
14	Planning time: 3.717 ms
15	Execution time: 3191286.075 ms

A Few Model Specific Databases

Relational DB



Graph DB



Search DB



Semi-structured DB



Timeserise DB



Spatial DB



Model Specific Databases

Key Value Storage



Data Processing Platform and dataframe technology



Flink 1.0




AWS RDS



DASK

And Many...



Why so many models and apps?

- In recent times, the database community has built many apps and techniques to handle various models and capabilities. (Like **relational**, **semi-structured**, or **networks** models and the capabilities like inverted index, data cube (Group by, CUBE by), centrality computation)
- These developments focused on a targeted vertical or solved domain-specific problems.
- Each of the apps developed and tuned best for its model and capabilities but incapable of the others.



- Relational Structure
- SQL
- Cube and Group queries
- Text search (Gin, GIST)
- Network queries
- Centrality computation




- Relational Structure
- Cypher
- Cube and Group queries
- Text search (Gin, GIST)
- Network queries
- Network analytics



- Text search
- Network queries
- Network analytics (Centrality, cluster, pagerank)

Polystore Database Systems



Design Goals of a Polystore Systems

- **Polystore should support location transparency** like federated databases (i.e., common query language).
- **Semantic Completeness:** The user will not lose any capabilities provided by its underlying storage engine.
- **Object Version Consistency:** The same version of the object should be available in multiple models.
- **Capability based Optimization :** Optimize the analytical computation depending on the capabilities.

Architectural Variations

Loosely Coupled

1. Cross model mediator based design, each provider will have a dedicated mediator to communicate with other providers.
2. Local storage has more control over the data, and the global controller maintains consistency and transparency.
3. Local operations are efficient, but model transformation cost is high.
4. Challenging to optimize analytical operations and create cross model materialized view and cross model index.

Tightly Coupled

1. Use a common interface to interact among stores, like a standard data frame or data structure for the whole polystore.
2. Local storage has less control over the data, and the central controller decides everything.
3. Transforming or rewriting queries from one store to another store is complex and ultimately boils down to a multi-query optimization problem.
4. It is hard to optimize the best plan for each store. The optimization cost is very high.
5. Easy to build a materialized view and index.

Hybrid



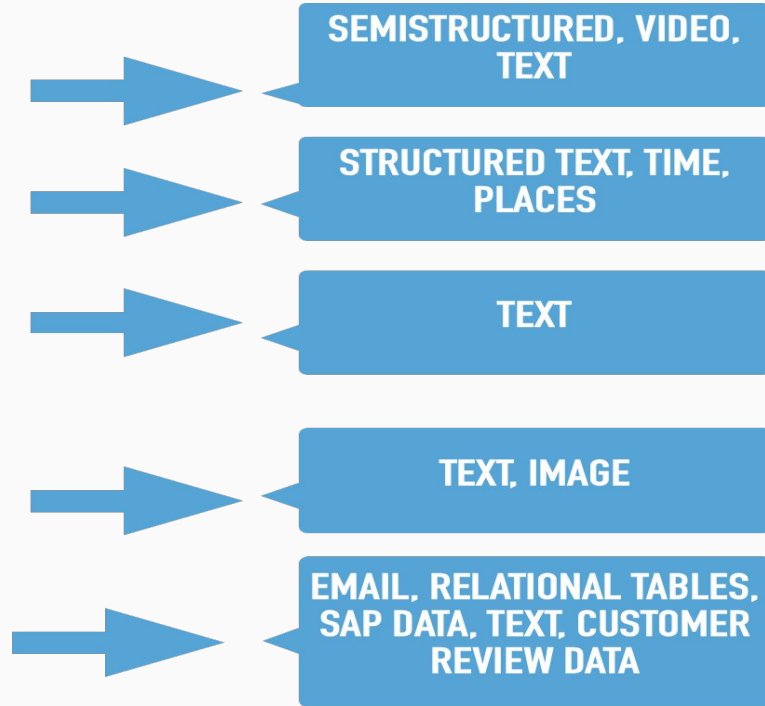
1. Trade-off between global control and local control
2. Very efficient for optimizing queries for each local storage.
3. Easy and efficient use of materialized view is possible.
4. Very much domain or vertical specific.

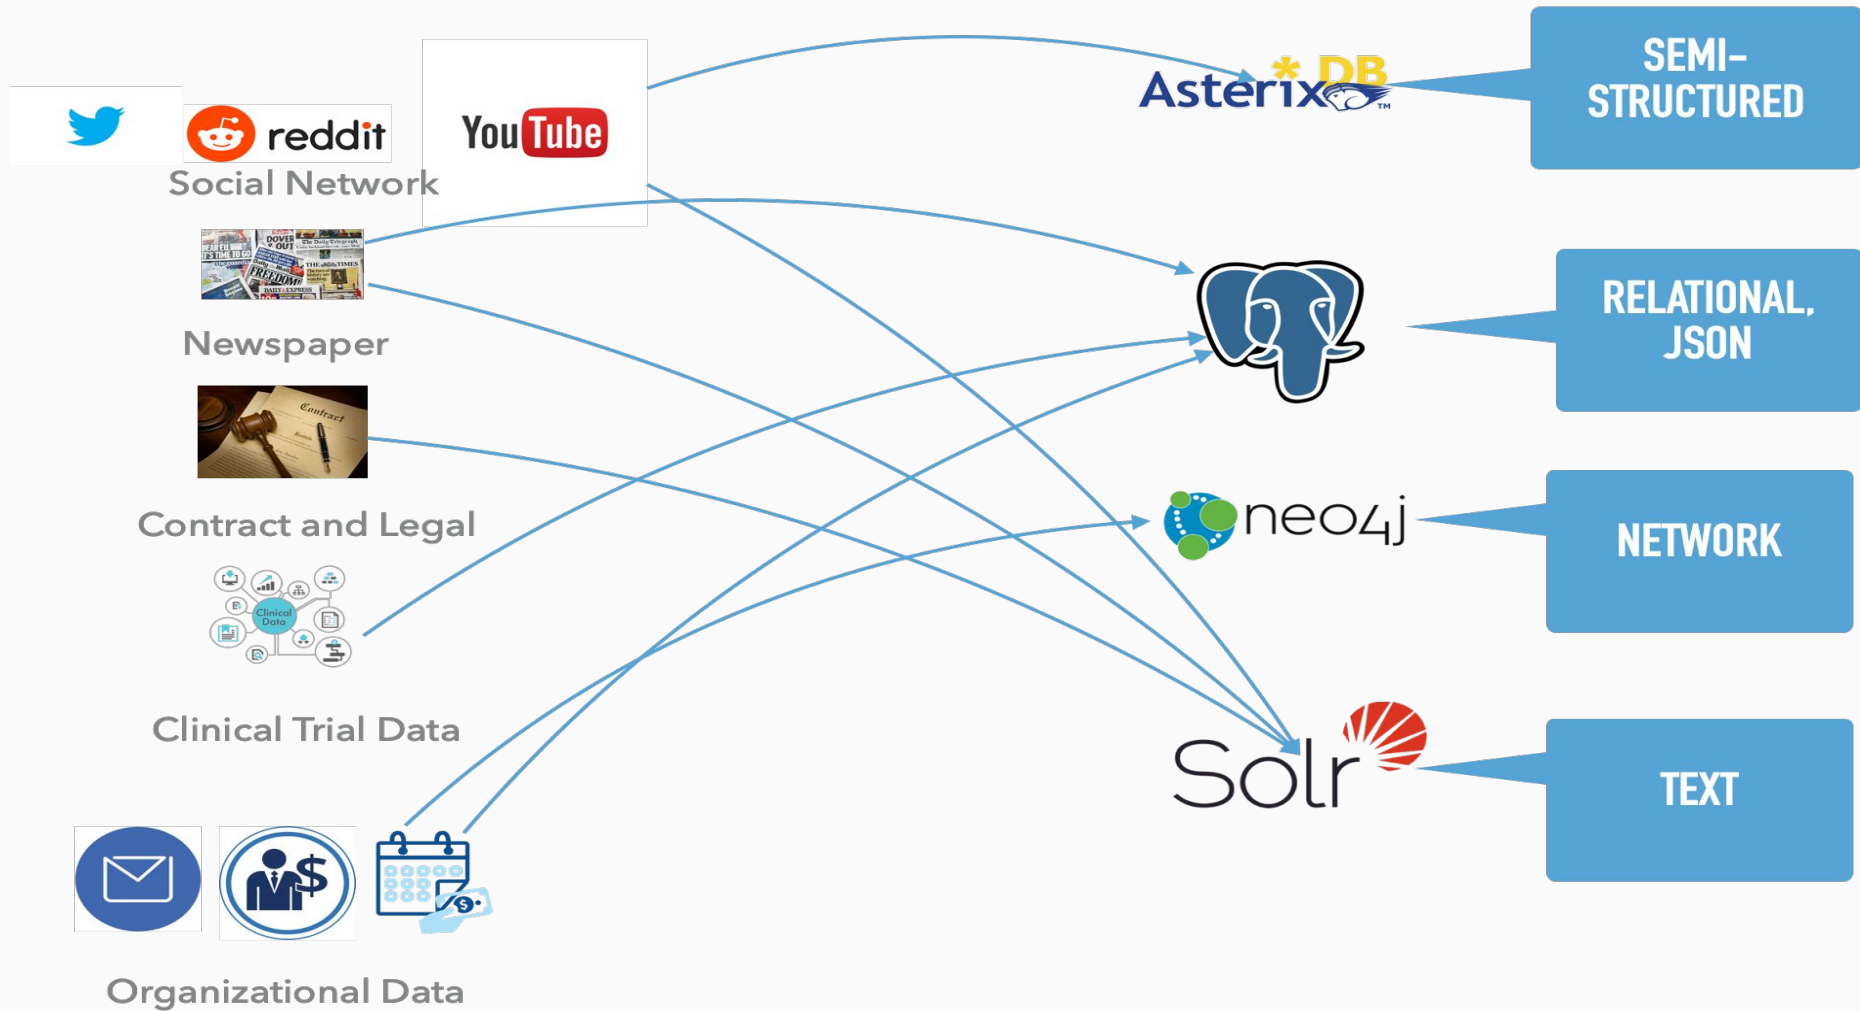
Example Polystore Systems

1. BigDAWG, MIT
2. CloudMdsQL, Inria
3. Estocada, UCSD and Inria
4. Polypheny-DB, University of Basel, Switzerland
5. Awesome, UCSD(*)
6. Polystore++ , Stanford University
7. Polybase, Microsoft.
8. OoX, HP-Labs

The Awesome Polystore

Data Variety





Social Sciences Questionnaire:

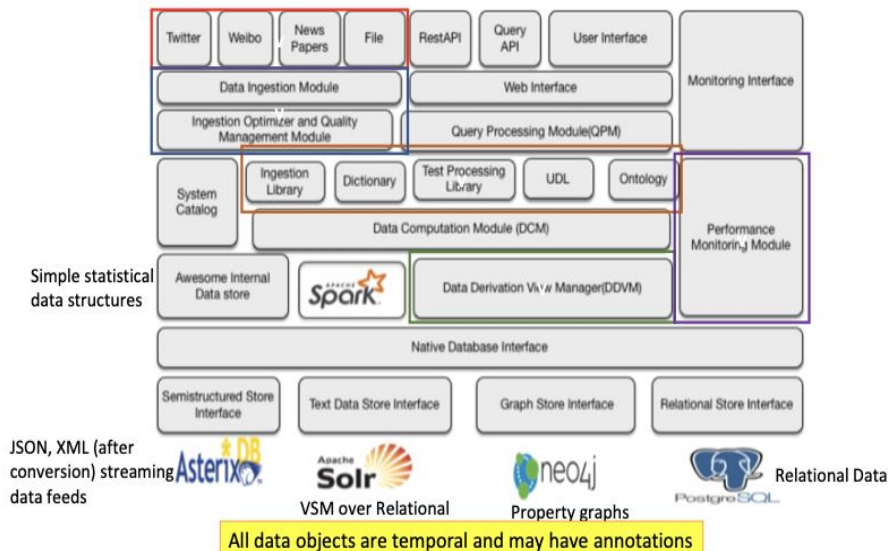
1. List all the accounts mentioned “Elizabeth Warren ” and also talking about “American DOS Movements.”
2. Find top 100 influential users those are co-spiking and talking about racial terms from the “Elizabeth Warren’s” network.
3. Find out top-k topics from the newspaper, those are also discussed in the “Elizabeth Warren’s” network.
4. Top k-topics discussed in the network but not covered by the newspaper.

Summary of Awesome Architecture

AWESOME integrates information over heterogeneous data

- A relational DBMS
- A graph DBMS
- A document/semi-structured DBMS
- A text search engine
- Vector and Matrix data from Analytics engines

AWESOME Polystore Architecture



The Knowledge Graph Case Study

Building a Knowledge Graph on the top of a polystore

The Problem

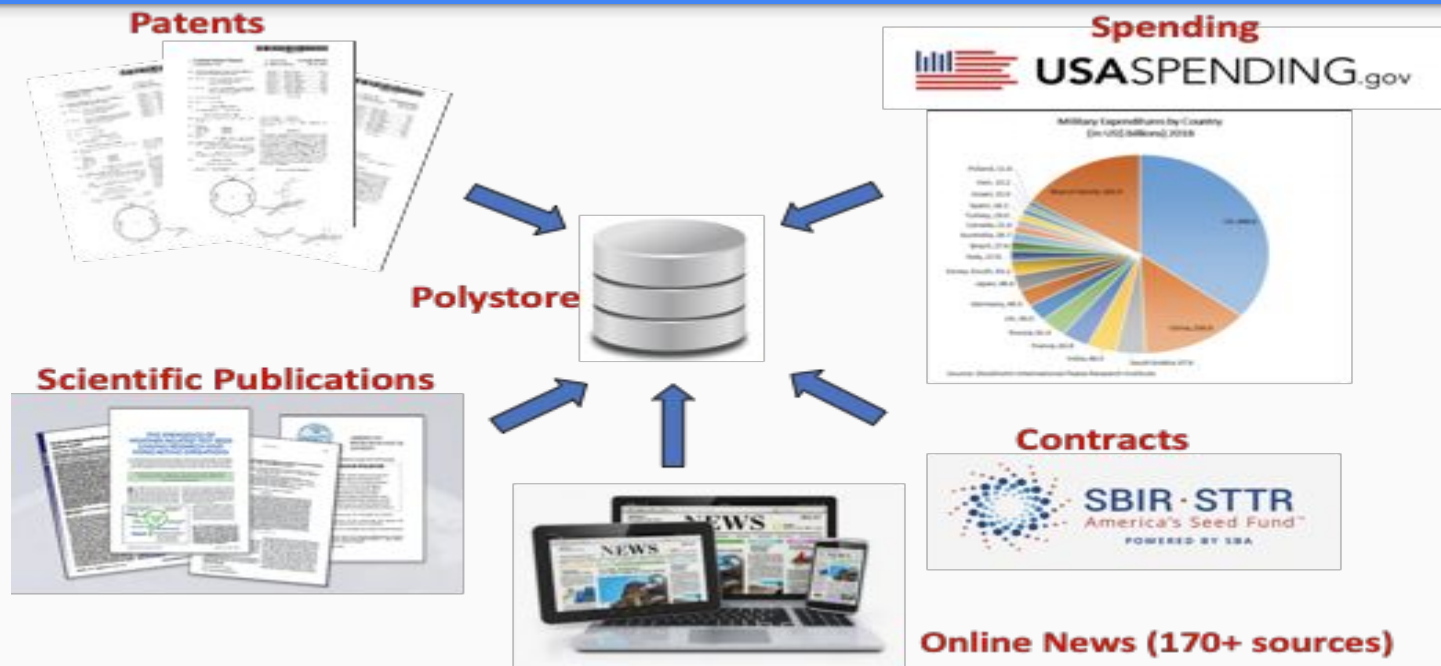
“Discover technology gaps in some domain and who can bridge the gaps?”

Data Set : all publicly available data

Solution Approach

- Create a Knowledge graph by assimilating information from multiple data sources.
- Search term expansion and association mining using KG
- Gap discovery using network query
- Potential partnership determination using Cube query

Building a Knowledge Graph Contd..



Query on Knowledge Graph

Search

1. Expand the search term automatically using the ontology.
2. Return a data cube with various distributions.

Gap Analysis

1. Find the technological associations by looking at network structure
2. Find the differences or gap of those structural associations.

Opportunity Analysis

Find collaboration opportunities from the network associations and other data.

Demo

Thanks!

Contact me:

sudasgupta@ucsd.edu

