


# SDSC Summer Institute 2021

## Scalable Machine Learning

Mai H. Nguyen & Paul Rodriguez  
2021-August-04



# Scalable Machine Learning Agenda

8:00 - 8:20 -- Machine Learning Overview

8:20 - 9:00 -- R on HPC

9:00 - 9:15 -- Break

9:15 - 10:15 -- Spark

**10:15 - 10:45 -- Spark Hands-On**

# Spark Hands-On

Mai H. Nguyen, Ph.D.

# SPARK PROGRAM STRUCTURE

- **Start Spark session**
  - `spark = SparkSession.builder.config(conf=conf).getOrCreate()`
- **Create distributed dataset**
  - `df = spark.read.csv("data.csv",header="True")`
- **Apply transformations**
  - `new_df = df.filter(col("dept") == "Sales")`
- **Perform actions**
  - `df.collect()`
- **Stop Spark session**
  - `spark.stop()`


# START SPARK SESSION

```
import pyspark
from pyspark.sql import SparkSession
```

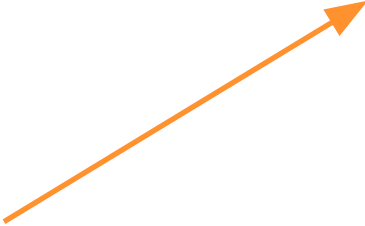
```
conf = pyspark.SparkConf().setAll([
    ('spark.master', 'local[*]'),
    ('spark.app.name', 'PySpark Demo')])
```

```
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```


Use \* to use all available cores, or integer value to specify number of cores to use



Configuration parameters for Spark session



Get existing Spark session or create new one



# LOAD DATA

- Loading data from local file system

```
df = spark.read.text("file:///<path>/<file>.txt")
```

```
df = spark.read.csv("file:///<path>/<file>.csv",  
                    header=True).cache()
```

- Loading data from HDFS

```
df = spark.read.text("hdfs:///<path>/<file>.txt") \  
    .cache()
```

```
df = spark.read.csv("hdfs:///<path>/<file>.csv",  
                    header=True).cache()
```

Indicates whether  
column headers exist



Cache data



# CHAINING

```
text_file = sc.textFile("hdfs://...")

counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)

counts.saveAsTextFile("hdfs://...")
```

**Chaining: Making multiple  
method calls on same object**



**Line continuation  
indicator**



## RDD Wordcount

# DROP ROWS WITH NULLS

- **Drop rows with null values**

```
df.dropna()  
df.dropna(how='any')  
df.dropna(how='all')
```

- **Check number of rows before and after dropping rows**

```
df.count()
```



# FILL IN MISSING VALUES

- **Replace null values with empty string**

```
df.na.fill('')
```

- **Count number of rows with nulls before and after filling nulls**

```
df.count()
```

# PARTITION DATA

- Partition available data into train and test data sets

```
train_df, test_df = df.randomSplit(0.8, 0.2), seed=<seed>)
```

Percentage of  
samples for train  
dataset



Percentage of  
samples for test  
dataset



# CREATE FEATURE VECTOR COLUMN

- **Create feature vector column**
  - Combines given list of columns into single vector column
  - To feed data to machine learning models

```
from pyspark.ml.feature import VectorAssembler
```

```
features = ['air_temp', 'relative_humidity']  
assembler = VectorAssembler(inputCols=features,  
                             outputCol='featureVector')
```

```
features_df = assembler.transform(df)  
features_df.show()
```

air_temp	relative_humidity
62.96	63.9



air_temp	relative_humidity	featureVector
62.96	63.9	[62.96, 63.9]

New column  
appended to  
features\_df

# SCALE DATA

- **Scale input data values**

- Standardize values to have zero mean and unit standard deviation
- Each feature is scaled separately
- Create scale transformer using train data, then apply to train/test data

```
from pyspark.ml.feature import StandardScaler
scaler = StandardScaler(inputCol="features_unscaled",
                        outputCol="features_scaled",
                        withStd=True, withMean=True)
scalerModel = scaler.fit(train_df)
```

```
train_df = scalerModel.transform(train_df)
test_df  = scalerModel.transform(test_df)
```

# BUILD MODEL


- **Build decision tree classifier**

- Create model
- Use fit() to train model

```
from pyspark.ml.classification import DecisionTreeClassifier
```

```
dt = DecisionTreeClassifier(  
    featuresCol='featureVector',  
    labelCol='label',  
    predictionCol='prediction')
```

Can specify name of  
columns for features,  
labels, and predictions



```
dt_model = dt.fit(<train>)
```

# APPLY MODEL

- **Apply trained model**
  - Use transform()

```
predictions = <model>.transform(<data>)
```

Will have <data> and  
new 'prediction'  
column appended at  
the end



Trained  
model

Input data

# EVALUATE CLASSIFICATION MODEL

- **Evaluator for classification model**
  - Calculates F1, precision, recall, accuracy

```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

mc_evaluator = MulticlassClassificationEvaluator(
    predictionCol='prediction',
    labelCol='label',
    metricName='f1')

mc_evaluator.evaluate(<predictions>)
```

Column with predictions

Column with labels

Evaluation metric

Contains predictions and labels

# PySpark Cluster Analysis Hands-On

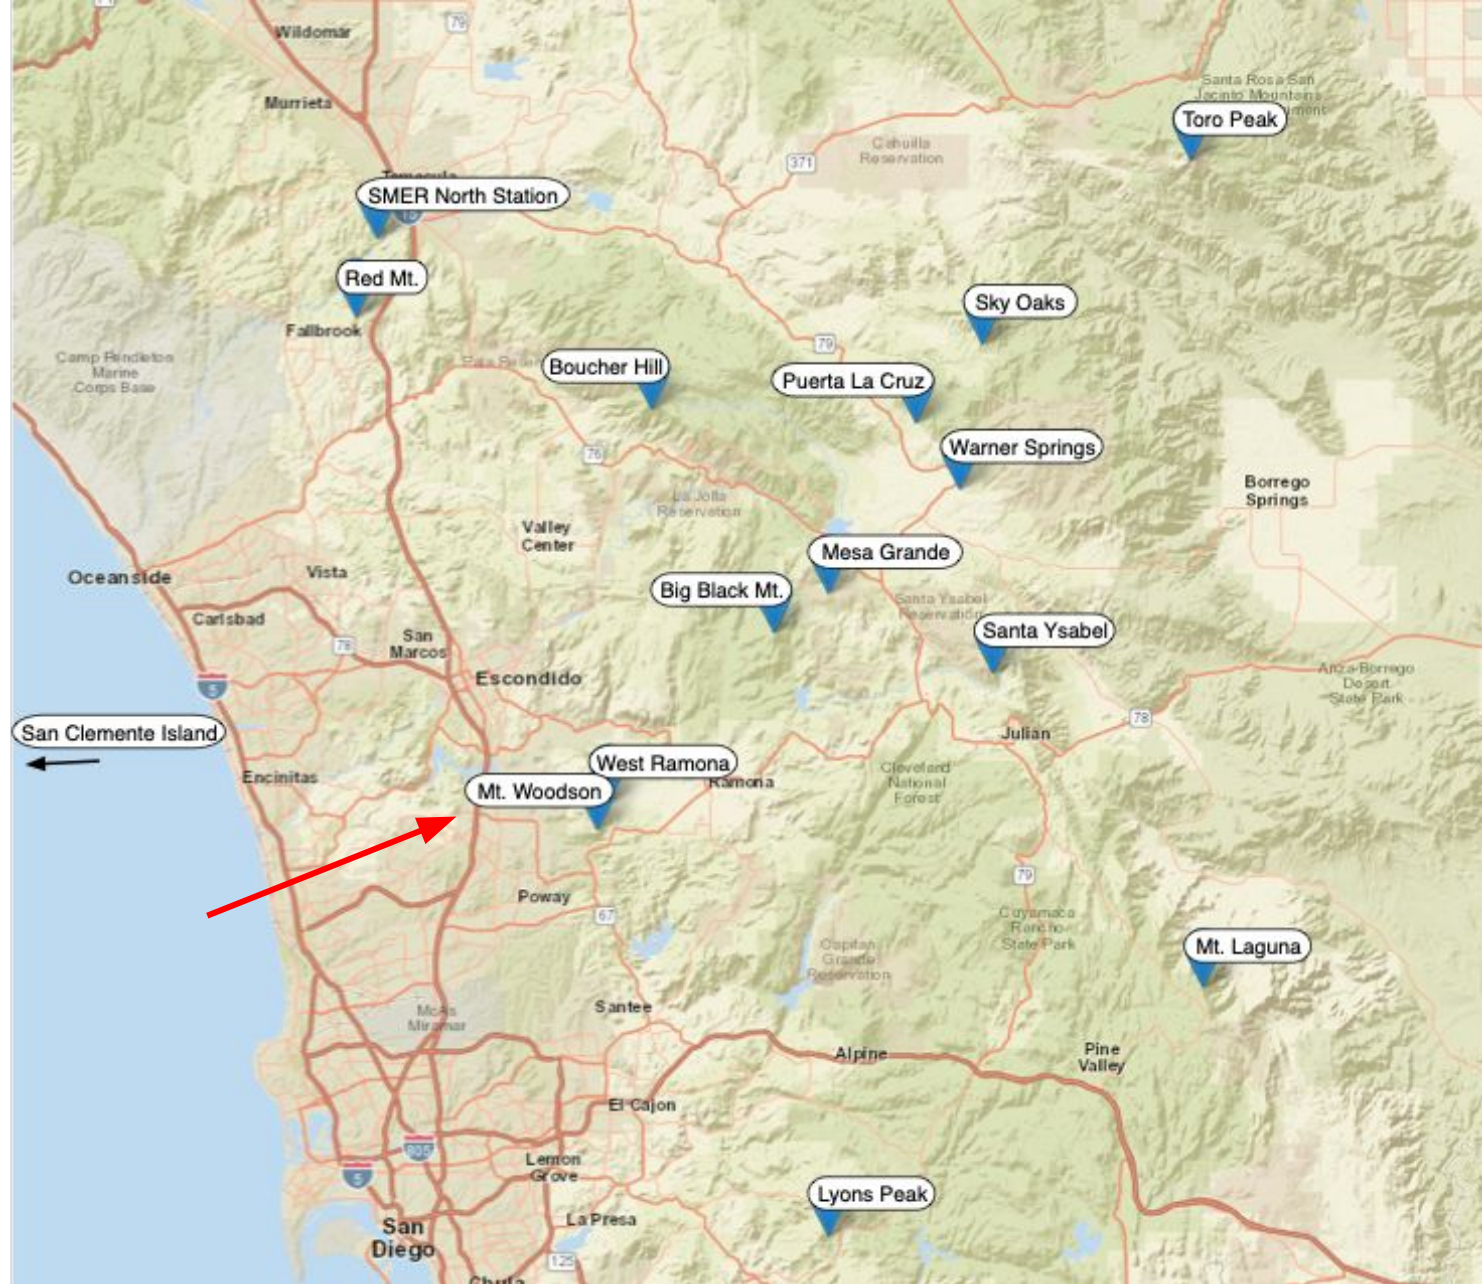
- **Data**
  - Weather station measurements
- **Task**
  - Perform cluster analysis to identify different weather patterns
- **Approach**
  - Spark k-means
- **Notebook**
  - `pyspark-clustering.ipynb`



# Dataset Description

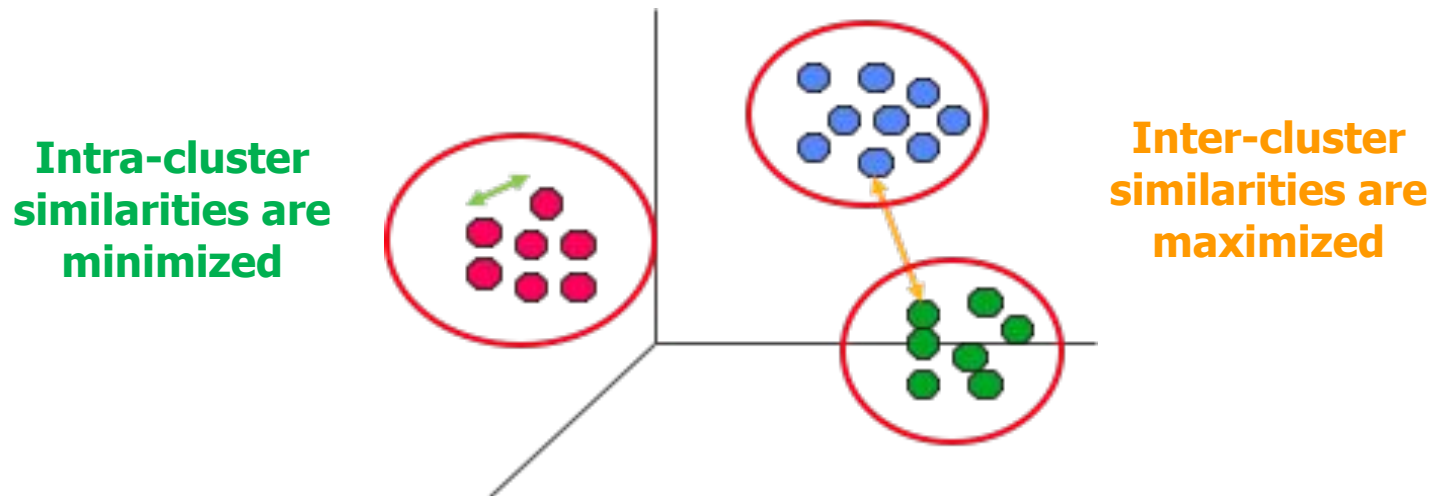
- **Measurements from weather station on Mt. Woodson, San Diego**
- **Air temperature, humidity, wind speed, wind direction, etc.**
- **Three years of data: Sep. 2011 - Sep. 2014**
  - minute\_weather.csv: measurement every minute
- **Source**
  - <http://hpwren.ucsd.edu>

# Map of HPWREN Weather Stations



# Cluster Analysis

- **Cluster analysis divides data into groups**
  - Grouping is based on some similarity measure.
  - Samples within a cluster are more similar to each other than to samples in other clusters.



<http://www-users.cs.umn.edu/~kumar/dmbook/index.php>

# *k*-Means Clustering

- **Partitional**
  - Clusters are divided into non-overlapping subsets
- **Centroid-Based**
  - Cluster represented by central vector
- **Simple, classic clustering technique**
  - Data points are grouped into  $k$  clusters
  - Cluster defined by cluster mean

- **Algorithm**

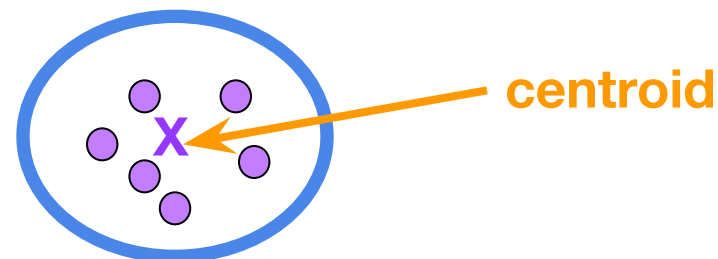
Select  $k$  initial *centroids* (cluster centers)

Repeat

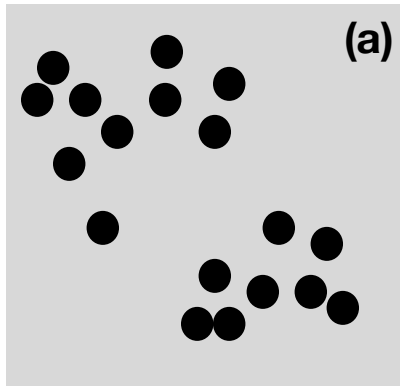
Assign each sample to closest centroid

Calculate mean of cluster to determine new centroid

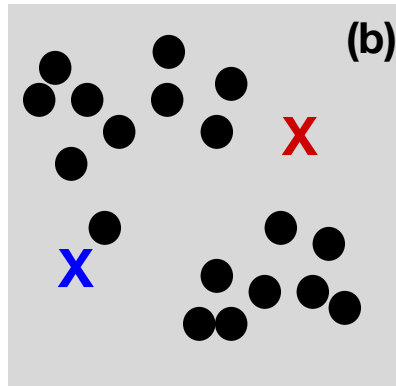
Until some stopping criterion is reached



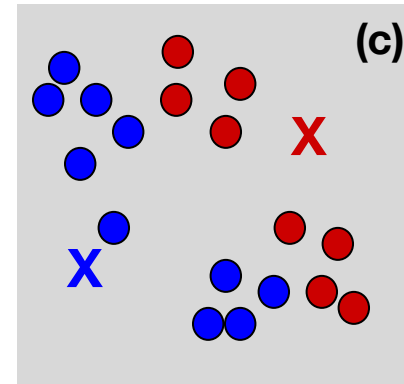
# *k*-Means Clustering Illustration



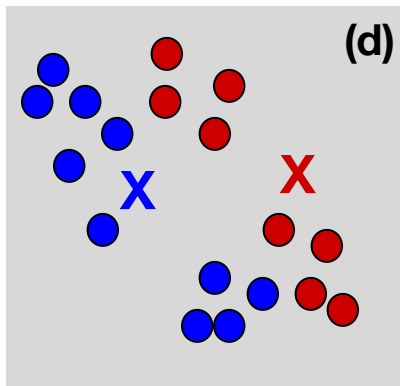
Original samples



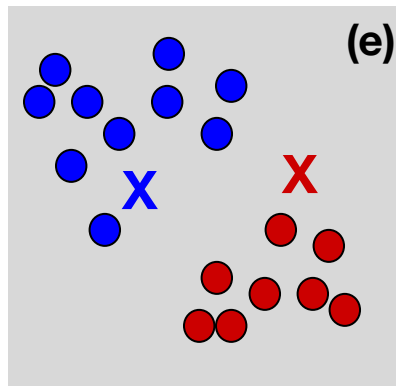
Initial Centroids



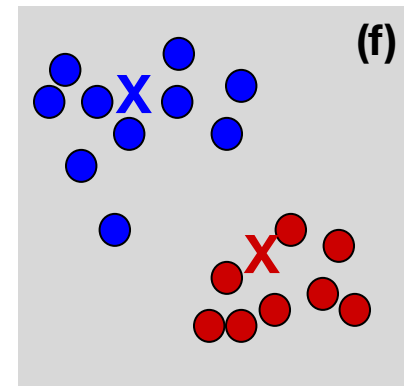
Assign Samples



Re-calculate Centroids




Assign Samples



Re-calculate  
Centroids

# Choosing Number of Clusters ( $k$ )

- **Elbow method**
  - Plot cluster evaluation metric (e.g., WSSE) vs. different values for  $k$
  - “Elbow” in plot suggests value(s) for  $k$



<http://stackoverflow.com/questions/6645895/calculating-the-percentage-of-variance-measure-for-k-means>

# Evaluating Clustering Results

- **Within-Cluster Sum of Squared Error (WSSE)**
- For each sample, error is distance to centroid.  
Then, **WSSE** is computed as:

$$WSSE = \sum_{i=1}^K \sum_{x \in C_i} \|x - m_i\|^2$$

$x$ : data sample in cluster  $C_i$

$m_i$ : cluster centroid (i.e., mean of cluster)

$\|x - m_i\|^2$ : Euclidean distance between  $m_i$  and  $x$



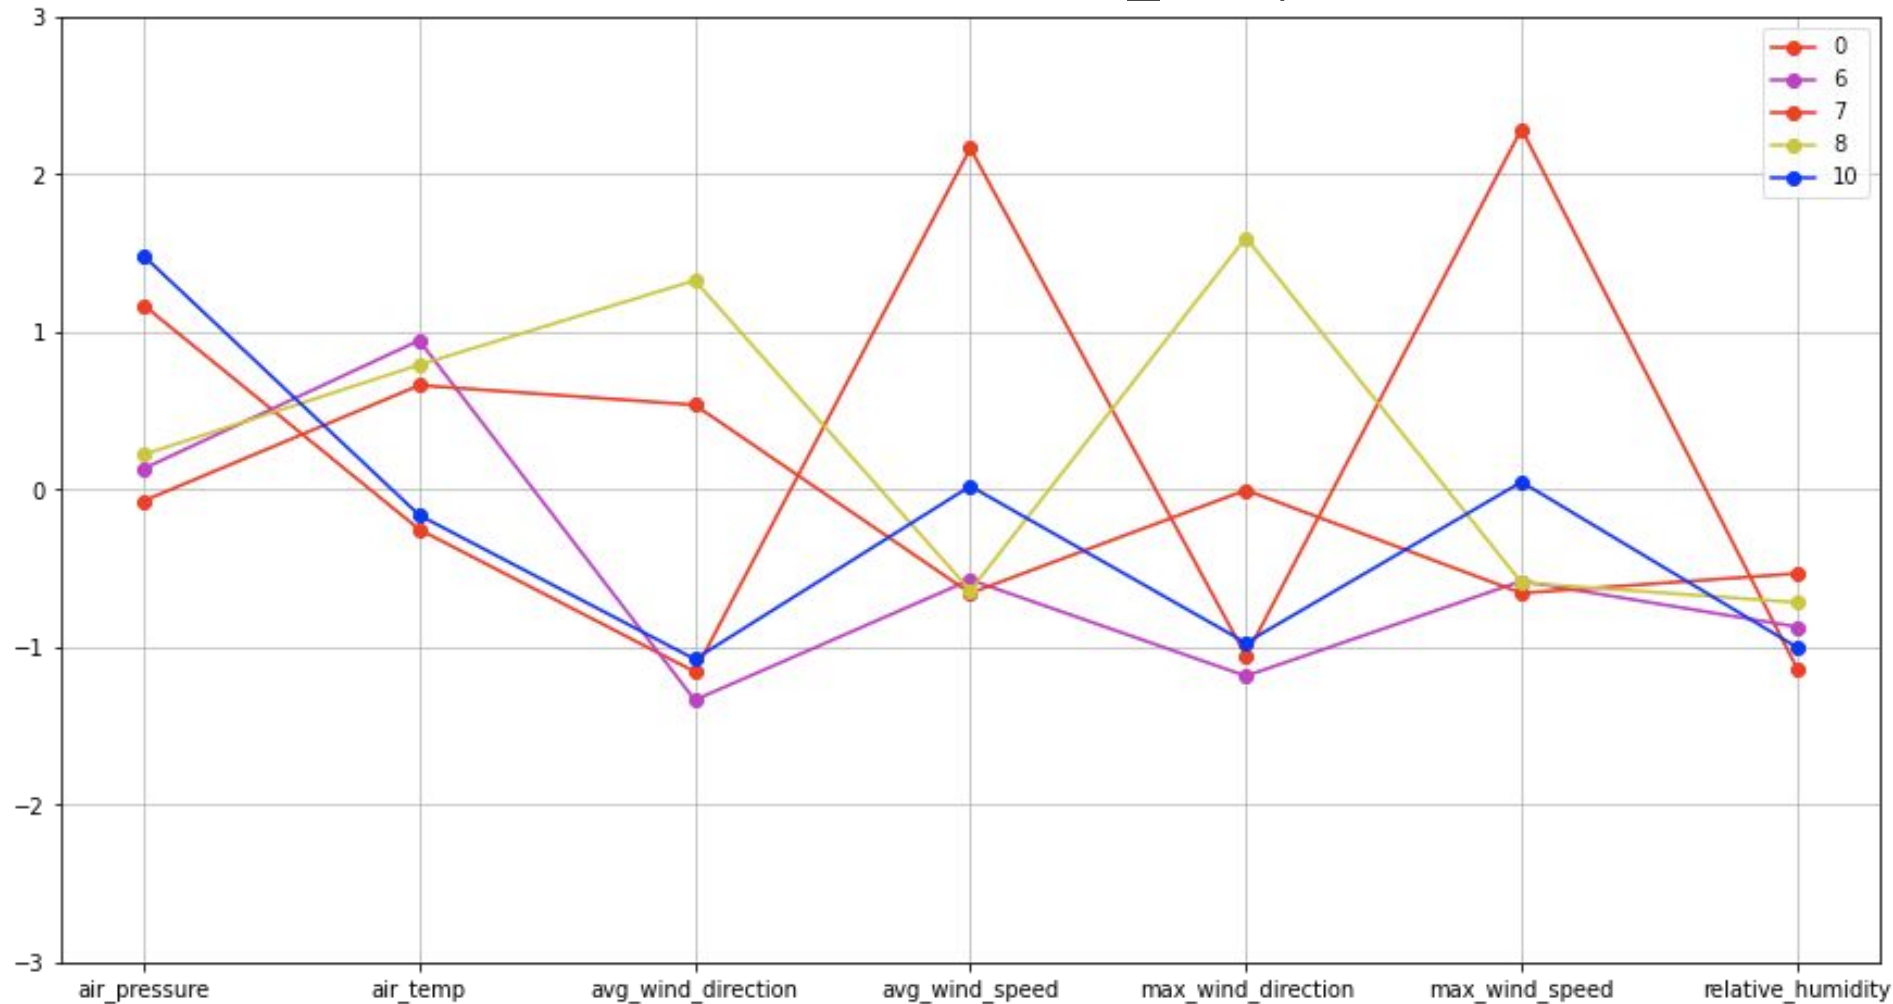
# Clustering Hands-On Overview

- **Setup**
  - Start Spark
  - Load modules
- **Load data**
  - Specify schema
  - Read in data from “minute\_weather.csv”
- **Explore data**
  - Look at schema, number of rows, summary statistics
- **Prepare data**
  - Drop nulls
  - Create feature vector
- **Perform k-means cluster analysis**
  - Use elbow plot to determine k
- **Evaluate clusters**
  - Plot cluster profiles
- **Stop Spark session**



# Cluster Profile: Parallel Plots

```
utils.parallel_plot(centersNamed[centersNamed['relative_humidity'] < -0.5],  
numClusters, colors=colors_used);
```



# Setup

- **Login to Expanse**

- Open terminal window on local machine
- `ssh login.expanse.sdsc.edu`

- **Pull latest from repo**

- `cd <your-SI-repo>`
- `git pull`
- URL: <https://github.com/sdsc/sdsc-summer-institute-2021>

# Server Setup for PySpark - Portal

- **Expanse Portal**

- <https://portal.expanse.sdsc.edu>
- Use trainXXX account
- Interactive Apps -> Jupyter

- **Parameters**

- Account: crl155
- Partition: shared
- Time limit (min): 60
- Number of cores: 2
- Memory required per node: 8 GB
- GPUs: 0
- Singularity image:  
/cm/shared/apps/containers/singularity/ciml/2021/pyspark-latest.sif
- Environment module: singularitypro
- Reservation: SI2021RES
- Type: JupyterLab

# Server Setup for PySpark - Command Line

- **In terminal window**
  - `cd 3.1b_Scalable_Machine_Learning`
  - `start-spark`
    - Alias for:
      - `export PATH="/cm/shared/apps/sdsc/galileo:${PATH}";`
      - `galileo launch --account crl155 --reservation SI2021RES --partition shared --cpus-per-task 2 --memory-per-node 8 --time-limit 01:00:00 --env-modules singularitypro --sif /cm/shared/apps/containers/singularity/ciml/2021/pyspark-latest.sif --bind /expance,/scratch,/cvmfs --quiet"`
  - Copy & paste URL in web browser
- **To check queue**
  - `squeue -u $USER`

# PySpark Cluster Analysis Hands-On

- **Code**

- **pyspark-cluster.ipynb**
  - Notebook for hands-on
  - Replace <<FILL-IN>> with code
- **pyspark-cluster-w-outputs.ipynb**
  - Has cell outputs
- **utils.py**
  - Has utility functions

- **Resources**

- [Apache Spark™ - Unified Analytics Engine for Big Data](#)
- [PySpark Documentation — PySpark 3.1.2 documentation](#)
- [Spark SQL and DataFrames - Spark 3.1.2 Documentation](#)
- Python for Data Science Cheat Sheet (pdf)