

SDSC Summer Institute 2021

Deep Learning

Mai H. Nguyen & Paul Rodriguez
2021-August-05



Deep Learning Agenda

- 8:00 - 8:45 -- Intro to Neural Networks / CNNs**
- 8:45 - 9:45 -- MNIST & TensorBoard Hands-On**
- 9:45 - 10:00 -- Break**
- 10:00 - 10:45 -- DL Layers & Architectures**
- 10:45 - 11:15 -- Lunch**
- 11:15 - 12:30 -- Transfer Learning Hands-On**
- 12:30 - 12:45 -- Break**
- 12:45 - 1:45 -- Deep Sequence Learning**
- 1:45 - 2:00 -- Wrap-Up**

Deep Learning Agenda

8:00 - 8:45 -- Intro to Neural Networks / CNNs

8:45 - 9:45 -- MNIST & TensorBoard Hands-On

9:45 - 10:00 -- Break

10:00 - 10:45 -- DL Layers & Architectures

10:45 - 11:15 -- Lunch

11:15 - 12:30 -- Transfer Learning Hands-On

12:30 - 12:45-- Break

12:45 - 1:45 -- Deep Sequence Learning

1:45 - 2:00 -- Wrap-Up

Introduction to Neural Networks and Convolution Networks (CNNs)

Paul Rodriguez, Ph.D.

Table of Contents

- **Overview of Neural Networks (aka Multilayer Perceptron)**
- **What is Deep Learning?**
- **Introduction to convolution and feature discovery**
- **Convolution Neural Networks**
- **MNIST exercise with Keras**

to get neural network:

Consider the Logistic Function

(aka sigmoid)

$$f(x_i) = \frac{1}{1 + \exp(-(b+wx))}$$


to get neural network:

Consider the Logistic Function

(aka sigmoid)


$$f(x_i) = \frac{1}{1 + \exp(-(b + wx))}$$

for parameters: $b = 0$, $w_1 = 1$




Make a graphical description of Logistic Function

$$f(x_i) = \frac{1}{1 + \exp(-(b+wx))}$$




Make a graphical description of Logistic Function

$$f(x_i) = \frac{1}{1 + \exp(-(b + wx))}$$



Make a graphical description of Logistic Function


$$f(x_i) = \frac{1}{1 + \exp(-(b + wx))}$$



this node (or unit) will transform input to output with logistic activation function

Make a graphical description of Logistic Function

$$f(x_i) = \frac{1}{1 + \exp(-(b + wx))}$$





this node (or unit) will transform input to output with logistic activation function



take derivatives of error (dE/dw) to update weights by gradient descent

How does changing parameters affect function?

$$f(x_i) = \frac{1}{1 + \exp(-(b + wx))}$$







How does changing parameters affect function?





Logistic function w/various weights

$$\text{for } y = 1 / (1 + \exp(-(b + w_1 * x)))$$




So combinations are highly flexible and nonlinear



(Note: these are both slightly shifted)

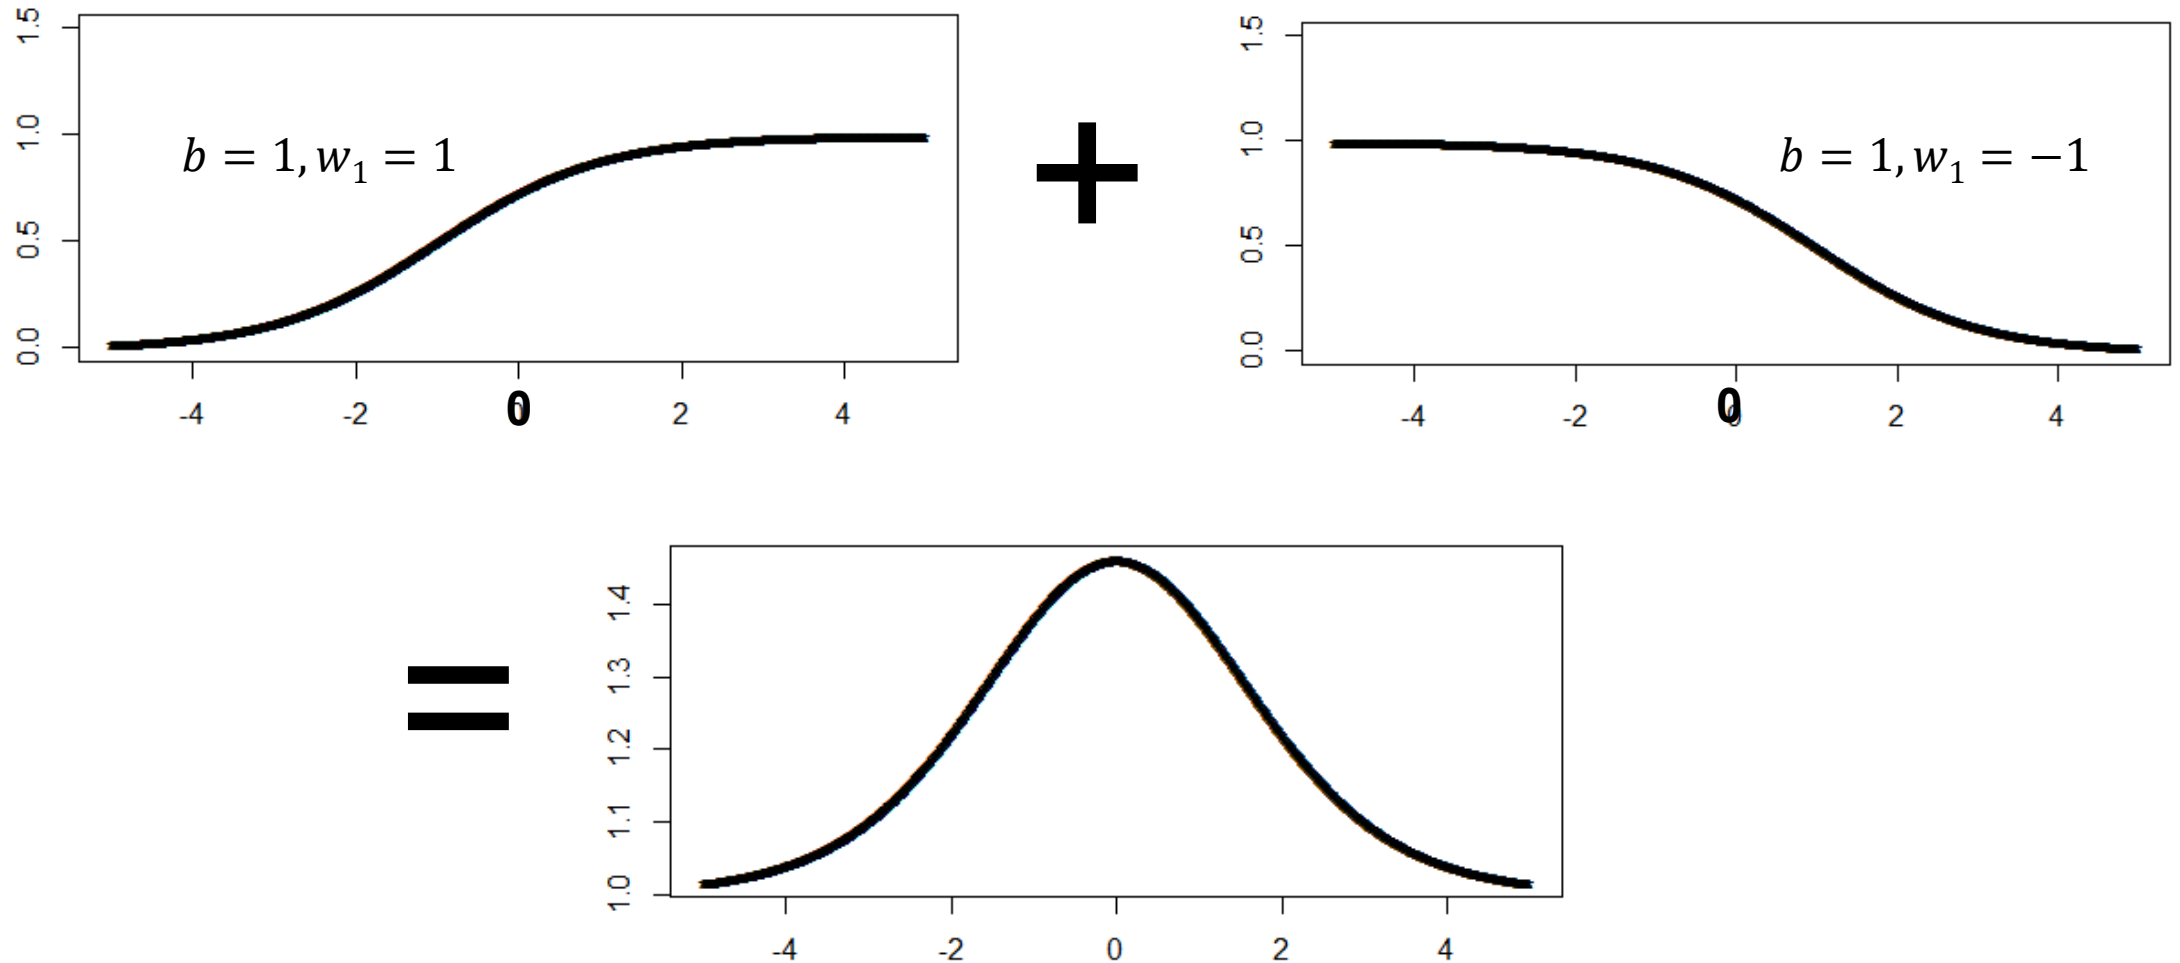
So combinations are highly flexible and nonlinear



If these are Hidden Layer Units,
and you add them....


then what does the output look like?

So combinations are highly flexible and nonlinear




Higher level function combinations


```
x=seq(-5,5,.1)
y1=1/(1+exp(10+ 10*x))
y2=1/(1+exp(-5+(-10)*x))
y3=1/(1+exp(1+1*y1+1*y2))
plot(x,y3,type="l")
```



```
y4=1/(1+exp(10+ (-10)*x))
y5=1/(1+exp(-5+(10)*x))
y6=1/(1+exp(1-1*y4-1*y5))
plot(x,y6,type="l")
```




```
y7=1/(1+exp(1+2*y3+1*y6))
plot(x,y7,type="l")
```




Higher level function combinations

```
x=seq(-5,5,.1)
y1=1/(1+exp(10+ 10*x))
y2=1/(1+exp(-5+(-10)*x))
y3=1/(1+exp(1+1*y1+1*y2))
plot(x,y3,type="l")
```




Multiple layer networks can represent any logical or real-valued functions (unbiased, but potential to overfit)


```
y4=1/(1+exp(10+ (-10)*x))
y5=1/(1+exp(-5+(10)*x))
y6=1/(1+exp(1-1*y4-1*y5))
plot(x,y6,type="l")
```



```
y7=1/(1+exp(1+2*y3+1*y6))
plot(x,y7,type="l")
```



we can add layers and nodes




we can add layers and nodes

Multilayer Perceptron

1 Input layer


1 Hidden layer

1 Output




we can add layers and nodes

Multilayer Perceptron




we can add layers and nodes

Multilayer Perceptron




Algorithm steps:



1. FORWARD PROPAGATE ACTIVATION:

apply input data x_i ,
calculate all node activations


algorithm steps:



1. FORWARD PROPAGATE ACTIVATION:
apply input data x_i ,
calculate all node activations

2. BACKWARD PROPAGATE ERROR:
calculate Error (or Loss) derivatives, dE/dY ,
pass it back to lower layer

algorithm steps:



1. FORWARD PROPAGATE ACTIVATION:


apply input data x_i ,
calculate all node activations

2. BACKWARD PROPAGATE ERROR:

calculate Error (or Loss) derivatives, dE/dY ,
pass it back to lower layer

For hidden layers use chain rule:
($dE/dY \quad dY/dH_3 \quad dH_3/dH_2$ etc...)
needs a summation of previous layer

algorithm steps:



1. FORWARD PROPAGATE ACTIVATION:

apply input data x_i ,
calculate all node activations


2. BACKWARD PROPAGATE ERROR:

calculate Error (or Loss) derivatives, dE/dY ,
pass it back to lower layer

For hidden layers use chain rule:
($dE/dY \quad dY/dH_3 \quad dH_3/dH_2$ etc...)
needs a summation of previous layers

***Beware: error signals get diluted as you go backward -
the 'vanishing gradient' problem***

algorithm steps:



**1. FORWARD PROPAGATE
ACTIVATION:**

apply input data x_i ,
calculate all node activations

2. BACKWARD PROPAGATE ERROR:
calculate Error (or Loss) derivatives (dE/dY)
pass it back to lower layer

3. Update weights and bias terms

$$w_{ji} = w_{ji} - \eta \frac{dE}{dw_{ji}}$$

NN Algorithm

NN Algorithm

INITIALIZE WEIGHTS (small random values)

NN Algorithm

INITIALIZE WEIGHTS

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

NN Algorithm

INITIALIZE WEIGHTS

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all derivatives to *minimize Loss (dL)*

NN Algorithm

INITIALIZE WEIGHTS

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all derivatives to *minimize Loss (dL)*

UPDATE WEIGHTS: $w \leftarrow w - \text{learning_rate} * \frac{dL}{dw}$

NN Algorithm

INITIALIZE WEIGHTS

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all derivatives to *minimize Loss (dL)*

UPDATE WEIGHTS: $w \leftarrow w - \text{learning_rate} * \frac{dL}{dw}$

STOP: when validation loss reaches minimum or converges

NN Algorithm

[heuristics, options to learn faster and/or better]

INITIALIZE WEIGHTS [use truncated distributions]

LOOP until stopping criterion: [work in batches of input]

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all derivatives to *minimize Loss (dL)*

UPDATE WEIGHTS: $w \leftarrow w - \text{learning_rate} * \frac{dL}{dw}$ [adapt learning rate, use momentum]


STOP: when validation loss reaches minimum or converges

[several metrics of loss are possible]

A heuristic for deep networks

RELU (rectified linear

activation function



$$f(a) = \begin{cases} a & a > 0 \\ 0 & a \leq 0 \end{cases}$$

where $a = XW$


A heuristic for deep networks

RELU (rectified linear

activation function

It is unscaled (bad!)

But df/da is constant (good!)



$$f(a) = \begin{cases} a & a > 0 \\ 0 & a \leq 0 \end{cases}$$

where $a = XW$


A heuristic for deep networks

RELU (rectified linear

activation function

It is unscaled (bad!)

But df/da is constant (good!)



$$f(a) = \begin{cases} a & a > 0 \\ 0 & a \leq 0 \end{cases}$$

where $a = XW$

RELU helps mitigates vanishing gradients

terminology and cheat sheet on output activations:

Type of Problem	Y outputs	Output Activation Function	Output PREDICTION (what you decide to predict)	Output Loss Function	Evaluative Measure
Regression: map into to real valued prediction	if $Y \in (-\infty, +\infty)^K$	(linear) $\hat{Y} = XW$	\hat{Y} :	Sum Squared Error (SSE)	Root Mean Squared Error (RMSE)
Multivariate output of 0's and 1's (multit-binary)	if $Y \in [0, 1]^K$	(sigmoid) $\hat{Y} = \frac{1}{1 + \exp^{-(XW)}}$	1 or 0	SSE	RMSE
Binary Classification	if $Y \in \{0, 1\}$	$\hat{Y} = \frac{1}{1 + \exp^{-(XW)}}$	A probability given by \hat{Y} : $P(y = 1 x)$	Cross Entropy $L = -y \log(\hat{y}) - (1 - y)(\log(\hat{y}))$	Accuracy, ROC
Multiclassification	if $Y \in \{0, 1\}^K$	$\hat{Y}_k = \frac{\exp^{-(XW_k)}}{\sum_k \exp^{-(XW_k)}}$	(softmax) Max class	Cross Entropy $L = - \sum_k y_k \log(\hat{y}_k)$	Accuracy

Some terminology and notes on output activations

Optimum is (often) at

$$\hat{y} = P(y = 1|x)$$

Type of Problem	Y outputs	Output Activation Function (this gives a SCORE \hat{Y} :)	Output PREDICTION (what you decide to predict)	Output Loss Function	Evaluative Measure
Regression: map into to real valued prediction	if $Y \in (-\infty, +\infty)^K$	$\hat{Y} = XW$	\hat{Y} :	Sum Squared Error (SSE)	Root Mean Squared Error (RMSE)
Multivariate output of 0's and 1's	if $Y \in [0, 1]^K$	$\hat{Y} = \frac{1}{1 + \exp^{-(XW)}}$	1 or 0	SSE	RMSE
Binary Classification	if $Y \in \{0, 1\}$	$\hat{Y} = \frac{1}{1 + \exp^{-(XW)}}$	A probability given by \hat{Y} : $P(y = 1 x)$	Cross Entropy $L = -y \log(\hat{y}) - (1 - y)(\log(\hat{y}))$	Accuracy, ROC
Multiclassification	if $Y \in \{0, 1\}^K$	$\hat{Y}_k = \frac{\exp^{-(XW_k)}}{\sum_k \exp^{-(XW_k)}}$	Max class	Cross Entropy $L = - \sum_k y_k \log(\hat{y}_k)$	Accuracy

Summary:

Pro:

Neural Networks in general, are flexible, powerful learners

Hidden layers learn a nonlinear transformation of input

Many heuristics about what works

Summary:

Pro:

- Neural Networks in general, are flexible, powerful learners
- Hidden layers learn a nonlinear transformation of input
- Many heuristics about what works

Con:

- Hard to interpret
- Needs more data
- Lots of parameters

What is deep learning?

What is deep learning?

Deep learning refers to learning complex and varied transformations of the input

What is deep learning?

Deep learning refers to learning complex and varied transformations of the input

Deep learning refers to **discovering** useful features of the input

What is deep learning?

Deep learning refers to learning complex and varied transformations of the input

Deep learning refers to **discovering** useful features of the input

Deep learning is a neural network with many layers

pause

onto Convolution Networks

Image features

- **MNIST - A database of handwritten printed digits**
(National Inst. of Standards and Technology)




Image features

- **MNIST - A database of handwritten printed digits**
(National Inst. of Standards and Technology)

How to classify digits?








Image features

- **MNIST - A database of handwritten printed digits**
(National Inst. of Standards and Technology)





How to classify digits?





Let's zoom into 5x6 window of pixels near the tip of '7'

Take a 3x3 patch of pixels and apply a 'filter' template – designed to find an edge





X

-1	0	+1
-1	0	+1
-1	0	+1

1. Multiply 3x3 patch of pixels with 3x3 filter

Let's zoom into 5x6 window of pixels near the tip of '7'

Take a 3x3 patch of pixels and apply a 'filter' template – designed to find an edge



(our weight parameters)



X

-1	0	+1
-1	0	+1
-1	0	+1

1. Multiply 3x3 patch of pixels with 3x3 filter “W”

Let's zoom into 5x6 window of pixels near the tip of '7'


Take a 3x3 patch of pixels and apply a 'filter' template – designed to find an edge





-1	0	+1
-1	0	+1
-1	0	+1

X

1. Multiply 3x3 patch of pixels with 3x3 filter “W”



2. Put answer in new cell of output map




-1	0	+1
-1	0	+1
-1	0	+1

X


1. Multiply 3x3 patch of pixels with 3x3 filter “W”

This is “image window” * W





2. Put answer in new cell of output map

3. Slide filter horizontally to get next output value



3	7		

3. Slide filter horizontally to get next output value



-1	0	+1
-1	0	+1
-1	0	+1



X

1. Multiply 3x3 patch of pixels with 3x3 filter "W"

2. Put answer in new cell of output map

3	7		

3. Slide filter horizontally to get next output value




-1	0	+1
-1	0	+1
-1	0	+1

1. Multiply 3x3 patch of pixels with 3x3 filter “W”

2. Put answer in new cell of output map


NOTE: sliding a filter is known as a “convolution” operation

3	7	3	



After vertical and horizontal sliding the 5x6 patch is now a 3x5 feature map.


3	7	3	-1
12	17	4	-1
15	21	4	-1



3	7	3	-1
12	17	4	-1
15	21	4	-1



After vertical and horizontal sliding the 5x6 patch is now a 3x5 feature map.

What do the highest values in the feature map represent?



3	7	3	-1
12	17	4	-1
15	21	4	-1

Optional next step:
Use another filter, and take maximum over elements -
“max pooling”





3	7	3	-1
12	17	4	1
15	21	4	-1

Optional next step:

Use another filter, and take maximum over elements - "max pooling"

2x2 filter has max=17

17		





3	7	3	-1
12	17	4	-1
15	21	4	-1

Optional next step:
Use another filter, and take maximum over elements -
“max pooling”

Slide filter ...

17	17	4
21	21	4



3	7	3	-1
12	17	4	-1
15	21	4	-1

After convolution and pooling 5x6 patch is **transformed** into a 2x3 feature map of 'edge gradients'

Slide filter ...

17	17	4
21	21	4

Feature engineering

In Computer Vision there are many kinds of edge detectors and many ways to scale them

-1	0	+1
-1	0	+1
-1	0	+1

But building features is hard, so if you have enough data ...

Convolution Neural Network (CNN)

In CNNs the filter values are weight parameters that are learned (**feature discovery**)

W_{11}	W_{12}	W_{13}
W_{21}	W_{22}	W_{23}
W_{31}	W_{32}	W_{33}

Convolution Neural Network (CNN)

In CNNs the filter values are weight parameters that are learned (**feature discovery**)

W_{11}	W_{12}	W_{13}
W_{21}	W_{22}	W_{23}
W_{31}	W_{32}	W_{33}

A convolution layer is a set of feature maps, where each map is derived from convolution of 1 filter with input

Convolution Neural Network (CNN)

More hyperparameters:

Size of filter (smaller is more general)

Convolution Neural Network (CNN)

More hyperparameters:

- Size of filter (smaller is more general)

- Number of pixels to slide over (1 or 2 is usually fine)

Convolution Neural Network (CNN)

More hyperparameters:

- Size of filter (smaller is more general)


- Number of pixels to slide over (1 or 2 is usually fine)

- Number of filters (depends on the problem!)

- Max pooling or not (usually some pooling layers)


Convolution with image

- Make 1 layer, using HxWx3 image (3 for Red,Green,Blue channels)




Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)




Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)




Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)



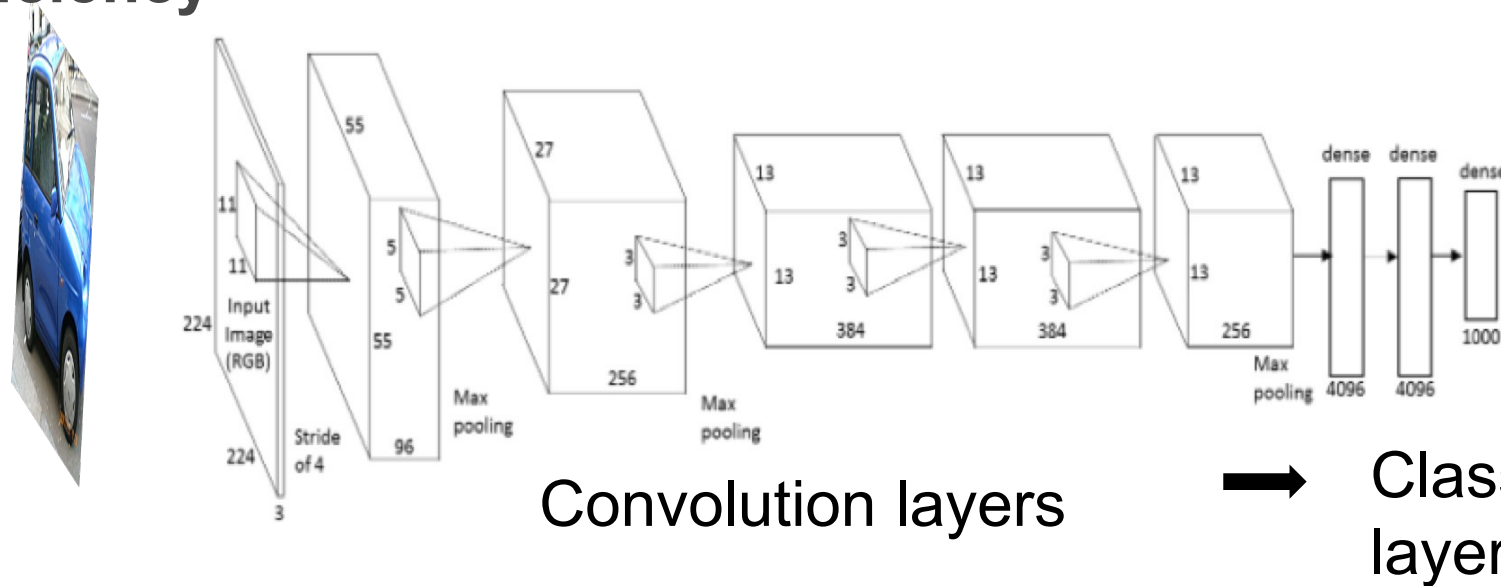
Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)



Large Scale Versions


- Large (deep) Convolution Networks are turning out to be feasible with GPUs (some are 100+ layers)
- Need large amounts of data and many heuristics to avoid overfitting and increase efficiency



What Learned Convolutions Look Like

First convolution layer filters are simple

features




What Learned Convolutions

First convolution layer filters are simple

features



Higher layers are more abstract features (or feature combinations)



Convolution Neural Network Summary

CNNs work because convolution layers have a special architecture and function – it is biased to do certain kind of transformations

Low layers have less filters that represent simple local features for all classes

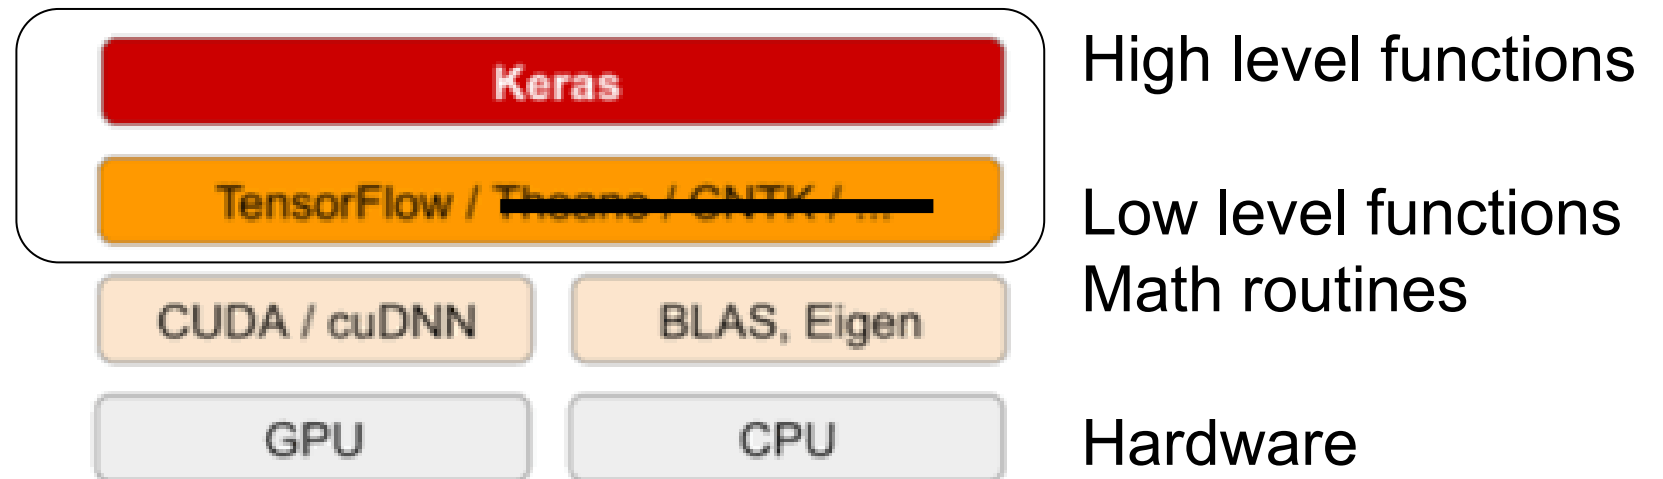
Higher layers have more filters that cover large regions that represent object class features

- **pause**

Deep Learning tools

“Tensorflow” handles vector and matrix multiplications, gradient calculations – on GPUs or CPUs

Keras (also from Google) is higher level library of functions to build networks – easiest to learn, widely used, well documented (now part of Tensorflow 2)



Keras commands build network in layers

```
mymodel = Sequential()
```

Define a model of sequential layers

Keras commands build network in layers

```
mymodel = Sequential()
```

Define a model of sequential layers

```
mymodel.add(Convolution2D ....
```

Add a convolution layer

Keras commands build network in layers

```
mymodel = Sequential()
```

Define a model of sequential layers

```
mymodel.add(Convolution2D (....
```

Add a convolution layer

```
mymodel.add(MaxPooling2D(....)
```

```
mymodel.add(Dropout(0.25))
```

Zero out nodes 25% of time for regularization

Keras commands build network in layers

```
mymodel = Sequential()
```

Define a model of sequential layers

```
mymodel.add(Convolution2D (....
```

Add a convolution layer

```
mymodel.add(MaxPooling2D(....)
```

```
mymodel.add(Dropout(0.25))
```

Zero out nodes 25% of time for regularization

```
mymodel.add(Dense( .....)
```


Add a fully connected layer

Tutorial: CNN for Digit Classification

- The ‘hello world’ of CNNs with Keras
- open SI2021_4.1b_Mnist_tutorial jupyter notebook

Note the “<<<< ----” marks places to change or fill in parameters

try different filter sizes
(as in following examples)



Expanse portal instructions

URL: `expanse.sdsc.edu`

Select User Portal ->

Select interactive session, Jupyter

Enter (others leave blank):

Account: `train####`

Partition: `compute`

Time: `120` (120 minutes, or whatever)

Cores: `128`

Memory: `248`

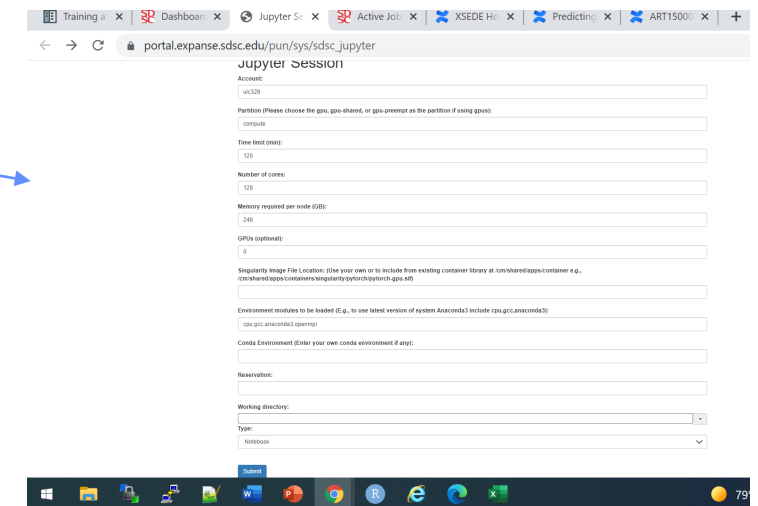
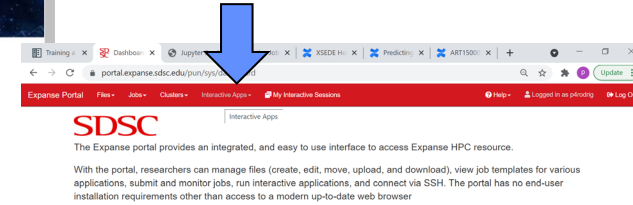
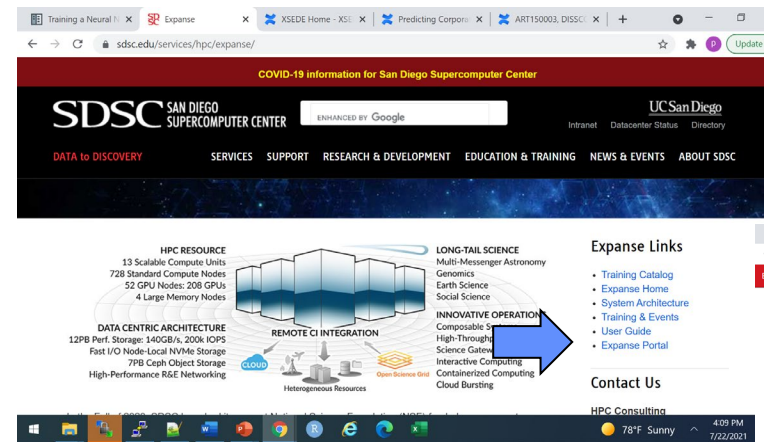
Singularity Image: `/cm/shared/apps/containers/
singularity/tensorflow/tensorflow-latest.sif`

Environment Modules: `singularitypro`

Reservation: `TBD`

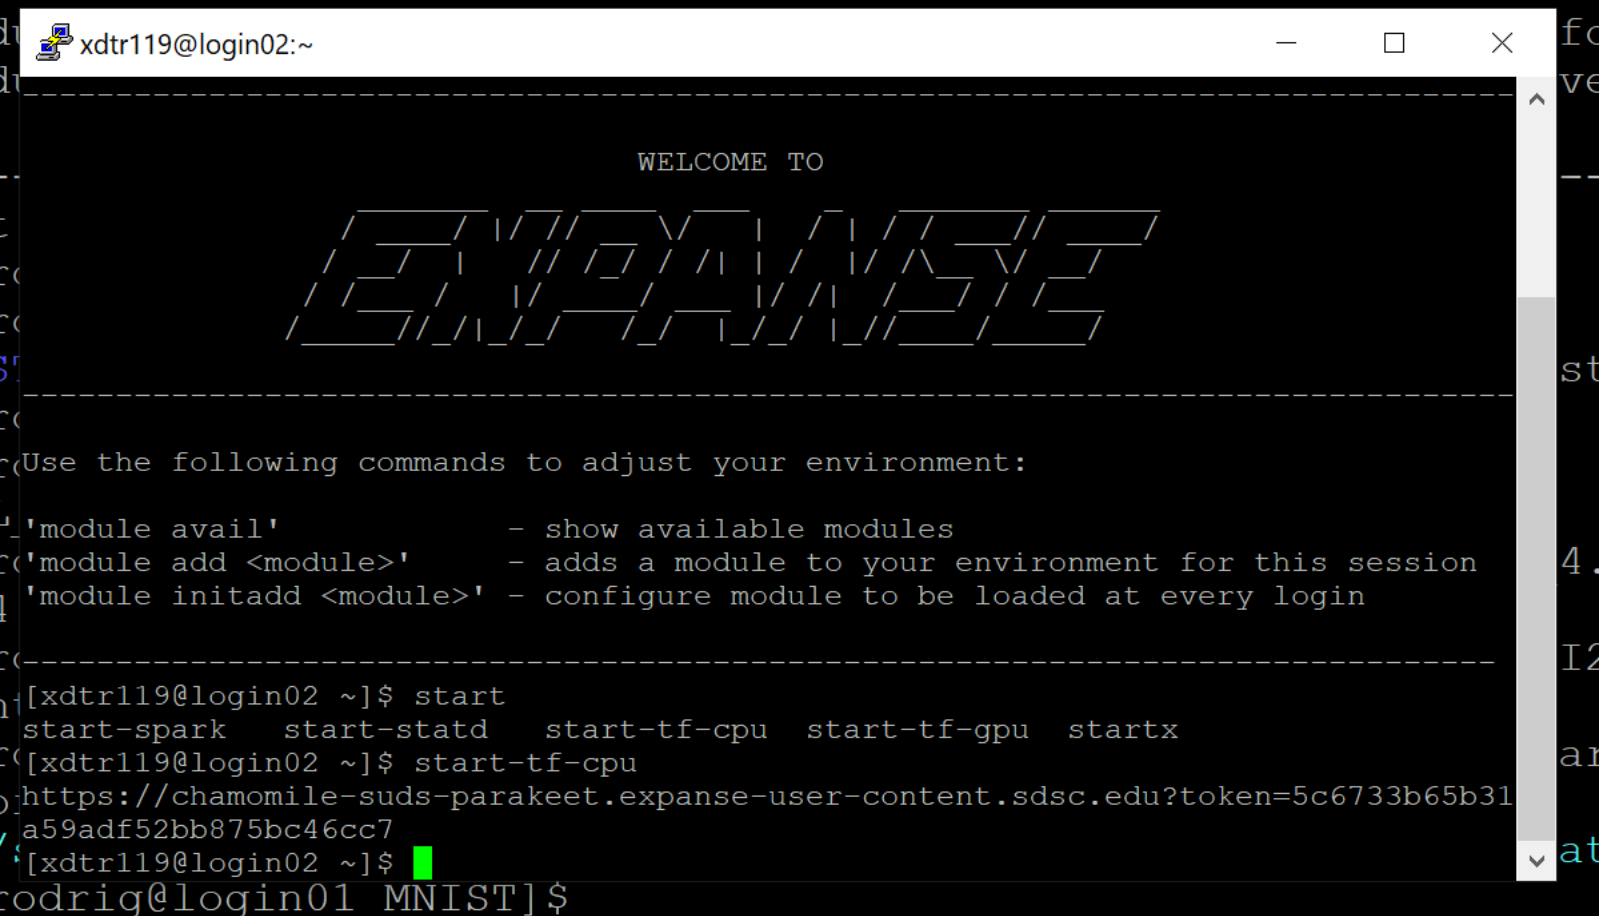
Working directory: `home`

Submit -> then select active Jupyter notebook job



Expanse terminal instructions

- 1
\$ start-tf-cpu
- 2
Cut & Paste URL to get notebook
- 3
Go to SI2021 4.1b directory and open:
SI2021_4.1b_Mnist_tutorial notebook

A terminal window titled 'xdtr119@login02:~' with a dark background. It displays a 'WELCOME TO' message followed by a large ASCII art logo for 'EXPANSE'. Below this, it lists commands to adjust the environment: 'module avail' (show available modules), 'module add <module>' (adds a module for this session), and 'module initadd <module>' (configures module to be loaded at every login). The user then runs 'start' and 'start-tf-cpu', which outputs a URL: 'https://chamomile-suds-parakeet.expanse-user-content.sdsc.edu?token=5c6733b65b31a59adf52bb875bc46cc7'. The prompt returns to the user, and another prompt 'rodrig@login01 MNIST]\$' is visible at the bottom.

```
xdtr119@login02:~  
-----  
WELCOME TO  
  
EXPANSE  
  
-----  
(Use the following commands to adjust your environment:  
'module avail' - show available modules  
'module add <module>' - adds a module to your environment for this session  
'module initadd <module>' - configure module to be loaded at every login  
-----  
[xdtr119@login02 ~]$ start  
start-spark start-statd start-tf-cpu start-tf-gpu startx  
[xdtr119@login02 ~]$ start-tf-cpu  
https://chamomile-suds-parakeet.expanse-user-content.sdsc.edu?token=5c6733b65b31a59adf52bb875bc46cc7  
[xdtr119@login02 ~]$  
rodrig@login01 MNIST]$
```


Training a Neural N ×

Expans ×

Dashboard - Expan ×

Jupyter Session ×

Netflix ×

+

⌵

📄

×

←

→

↻

portal.expans.sdsc.edu/pun/sys/sdsc_jupyter

🔍

☆

⚙

☰

p

Update ⋮

Open OnDemand / Jupyter Session

Jupyter Session

Account:

sds184

Partition (Please choose the gpu, gpu-shared, or gpu-preempt as the partition if using c

compute

Time limit (min):

90

Number of cores:

128

Memory required per node (GB):

248

GPUs (optional):

0

Singularity Image File Location: (Use your own or to include from existing container lib



Training a Neural N ×

Expans ×

Dashboard - Expan ×

Jupyter Session ×

Netflix ×

+

⌵

📄

×

←

→

↻

portal.expans.sdsc.edu/pun/sys/sdsc_jupyter

🔍

☆

⚙

☰

p

Update ⋮

Singularity Image File Location: (Use your own or to include from existing container library at /cm/shared/apps/container e.g., /cm/shared/apps/containers/singularity/pytorch/pytorch-latest.sif)

/cm/shared/apps/containers/singularity/tensorflow/tensorflow-latest.sif

Environment modules to be loaded (E.g., to use latest version of system Anaconda3 include cpu,gcc,anaconda3):

singularitypro

Conda Environment (Enter your own conda environment if any):

Reservation:

QoS:

Working directory:

home

Type:


JupyterLab

Submit

Windows taskbar with icons for File Explorer, Edge, Word, PowerPoint, JupyterLab, and a green circular icon.

70°F Clear 8:56 PM 8/1/2021

3x3 first convolution layer filter and activation



9x9 first convolution layer filter and activation

