

SDSC Summer Institute 2021

Running Secure Jupyter Notebooks on Expanse

Preparation Day – Wednesday, July 28

Mary Thomas
SDSC

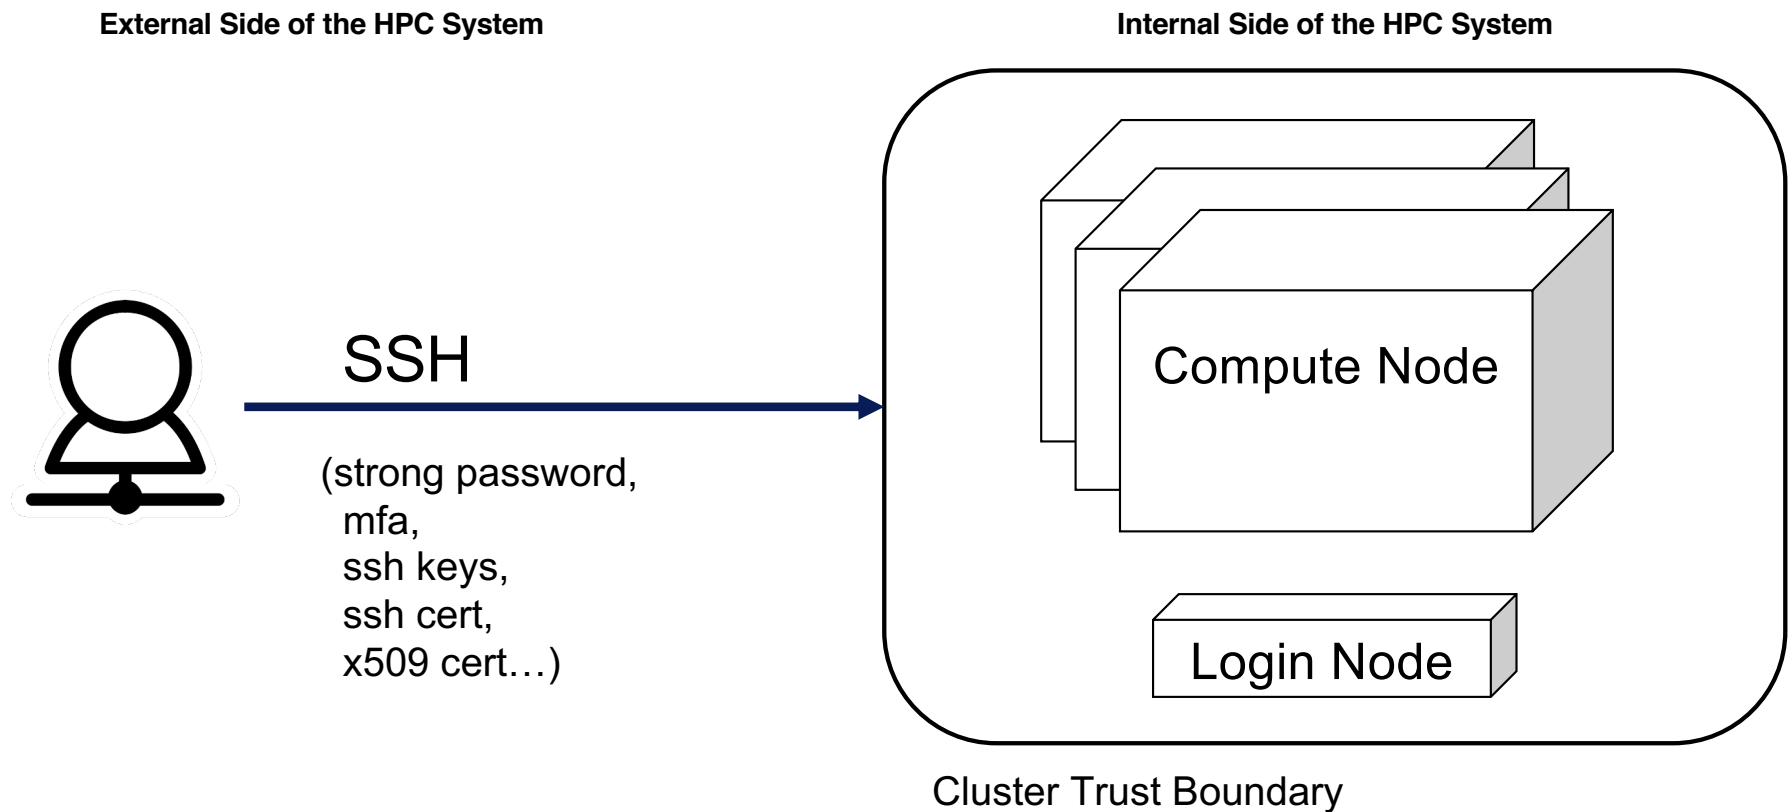
EXPANSE
COMPUTING WITHOUT BOUNDARIES

SAN DIEGO SUPERCOMPUTER CENTER

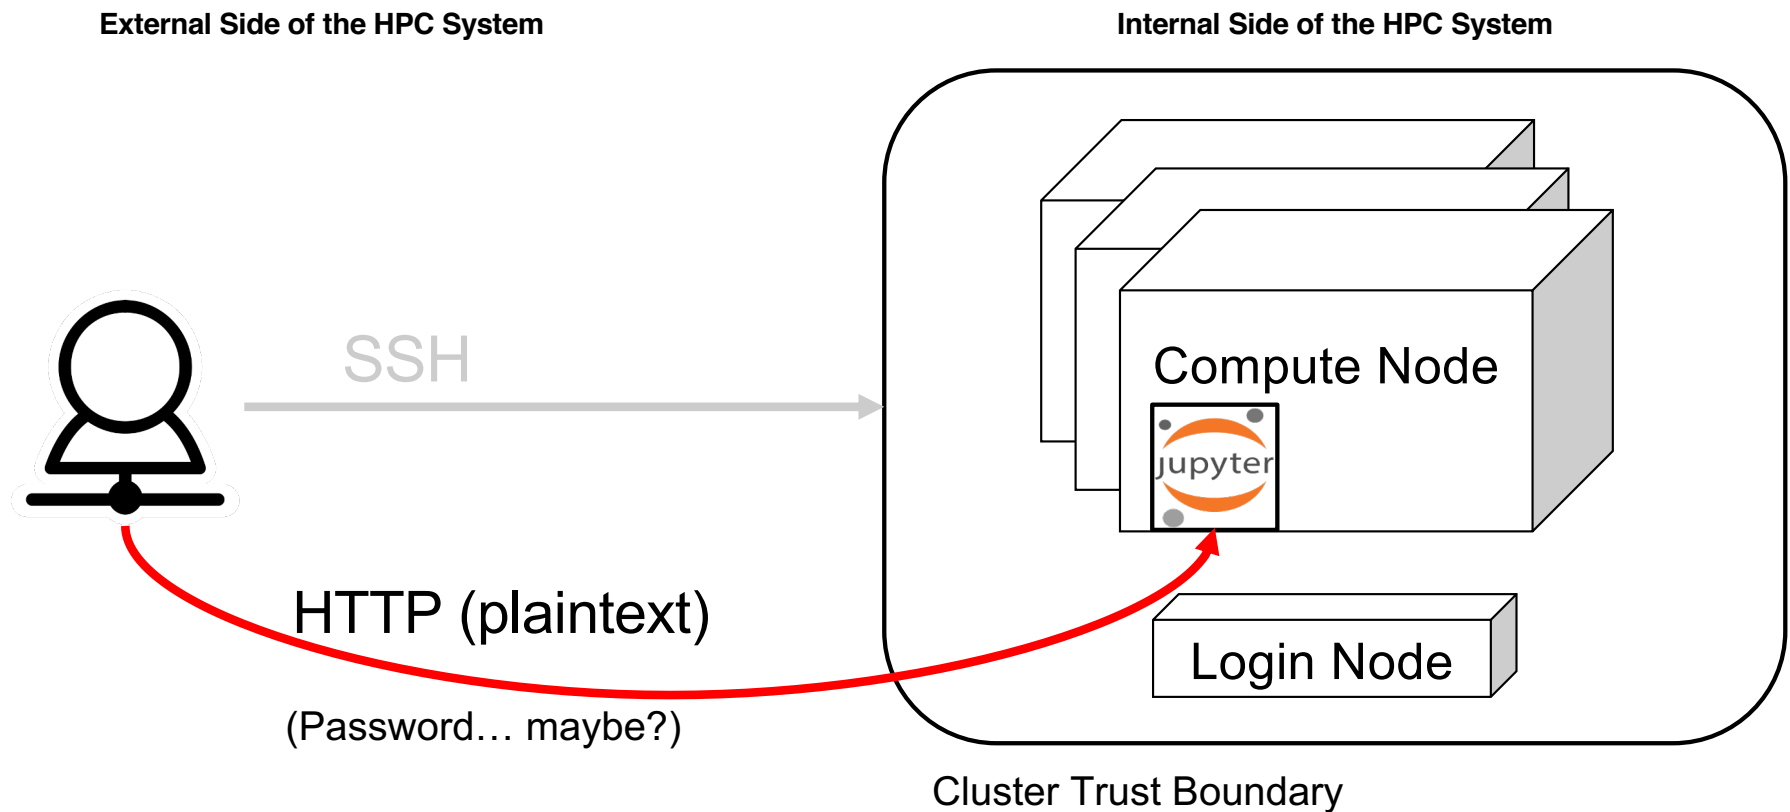


NSF Award 1928224

Secure Connections to HPC Systems are Important



Jupyter Notebooks are Popular



But provide a plaintext back-door to the system

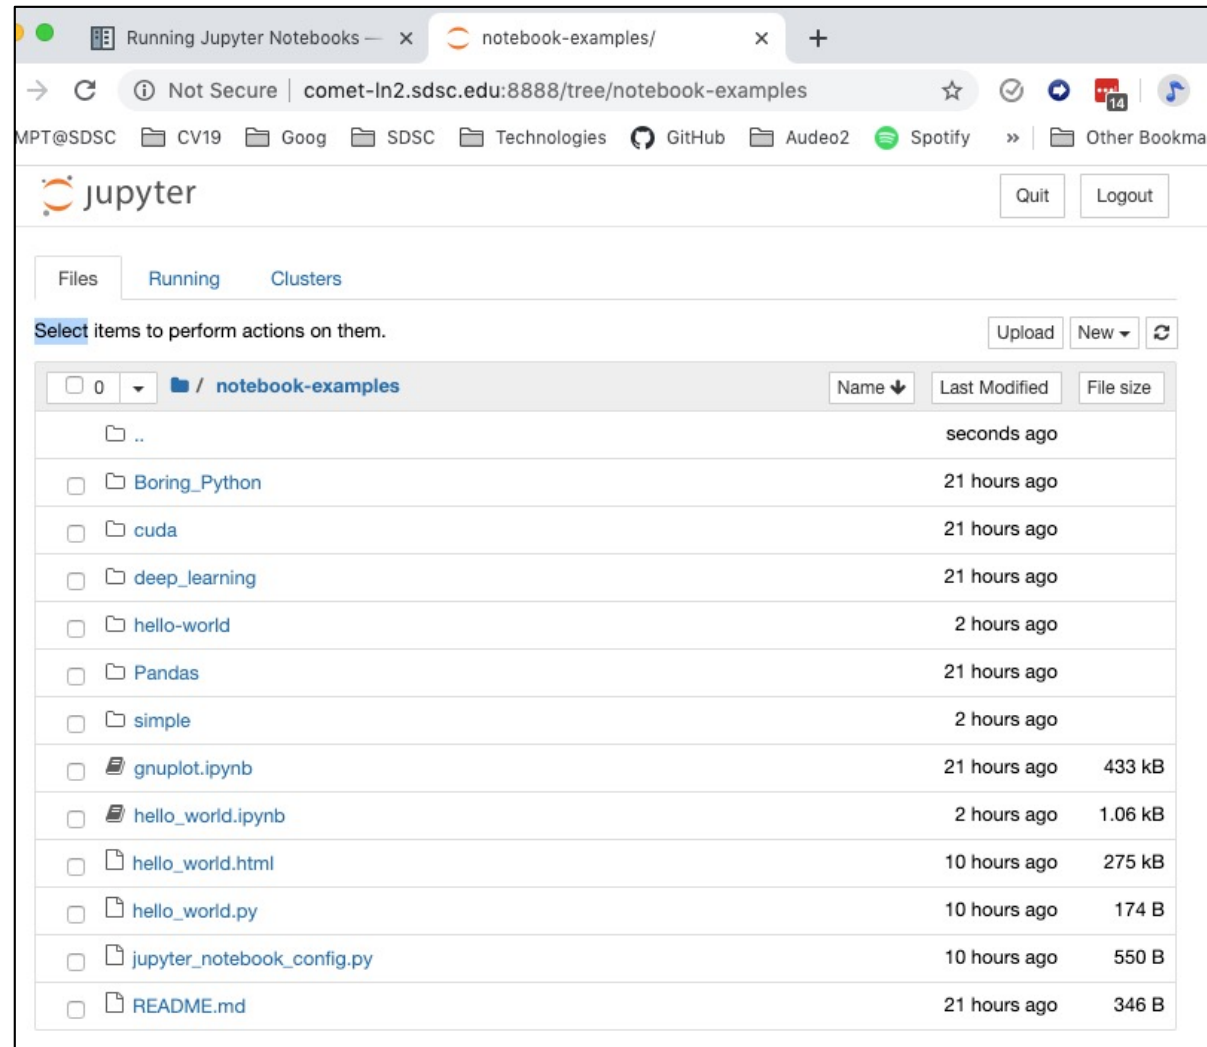
Jupyter Has a Key Vulnerability: Jupyter Servers Provide Access to HPC File Systems

SDSC Jupyter Services Policy:

- Portals, JupyterHub, and other services cannot be mounted directly to disk (must be on VM or external)
 - Many use root in vulnerable ways
 - If a user launches Jupyter Lab or Notebooks, the jobs will be killed.
- No applications can run on login nodes

SDSC recommendation:

- use secure connections:
when you choose insecure connections your account is vulnerable to hacking



Motivation: Make Doing the Right Thing Easier than the Wrong Things

A Wrong Thing: Plaintext to Compute Node

- Submit batch job.
- Wait till job runs.
- Figure out what node it's on.
- Point web browser at node.

A Right Thing: Improve secure access:

- Invoke **the Satellite Reverse Proxy Service**
- Point browser at secure, encrypted URL (HTTPS).
- (Wait until Jupyter Notebook shows up.)

A Simple Solution: SDSC Satellite Reverse Proxy Service

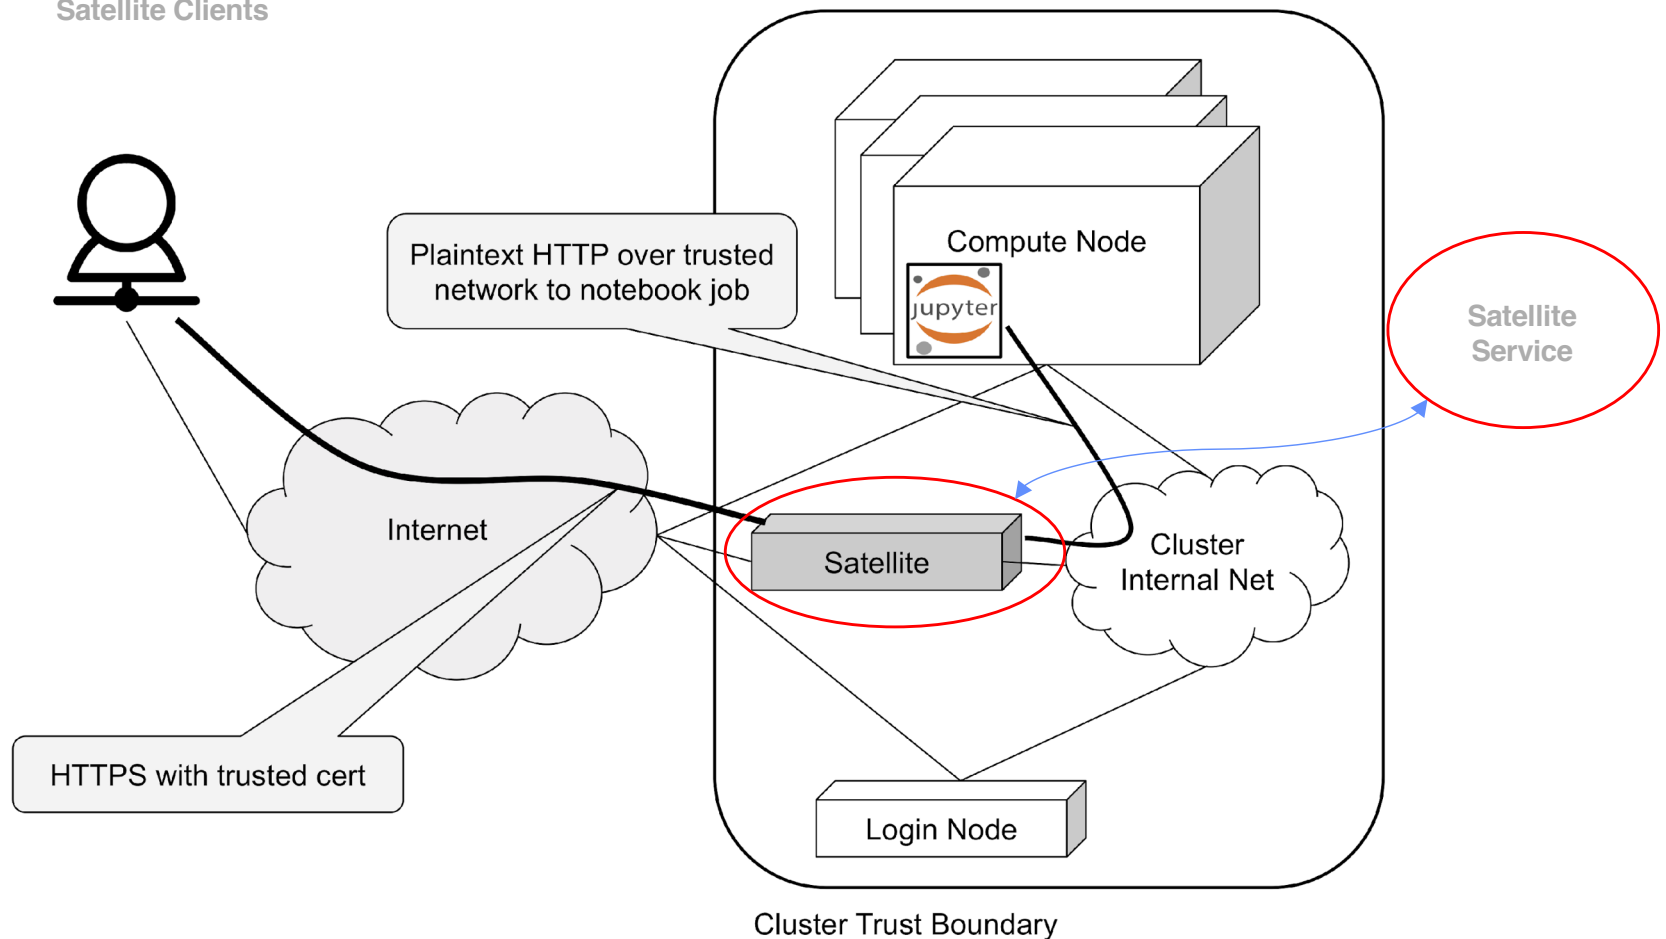
Just Two Components!

- **Satellite:** a self-service HTTP(s) reverse-proxy.
- **Satellite Client:** a shell-based utility to orchestrate a user's interaction with both Satellite and Slurm to start a Jupyter session within a batch job.

Jupyter Notebooks With Satellite

External Side of the HPC System:
Satellite Clients

Internal Side of the HPC System



Satellite Clients

start-jupyter

- 1st generation shell utility developed to orchestrate a user's interaction with both Satellite and Slurm to start a Jupyter session within a batch job.
- Key features in design:
 - User calls start-jupyter launch script, which requests token from Satellite, passes token to batch job script and submits the job to Slurm; token redeemed from batch job once it runs
 - Provided user with a prefabricated set of batch job scripts to choose from for certain popular applications on each system; user could modify/make their own custom batch job script
 - Small custom shell function library to make code more reusable
- Currently runs on: Expanse, Comet, TSCC, TSCC Stratus

<https://github.com/sdsc-hpc-training-org/reverse-proxy>

galileo

- 2nd generation shell utility developed to orchestrate a user's interaction with both Satellite and Slurm to start a Jupyter session within a batch job.
- Developed while reviewing start-jupyter codebase to sort out how best to support Expanse (OOD) Portal and HPC User Services Group long-term; effectively recycled existing an SSH tunneling orchestration utility to use Satellite proxy service instead.
- Key features in design:
 - Recreate same interactions with Satellite service.
 - Increase flexibility for users to configure software environment; but also try to make it simpler for them to do themselves
 - Batch job script is generated completely on-the-fly.
 - Command-line argument driven.
 - Quiet mode for OOD portal

<https://github.com/mkandes/galileo>

galileo demo examples on Expanse

```
# Set location of galileo directory on Expanse
```

```
export PATH="/cm/shared/apps/sdsc/galileo:${PATH}"
```

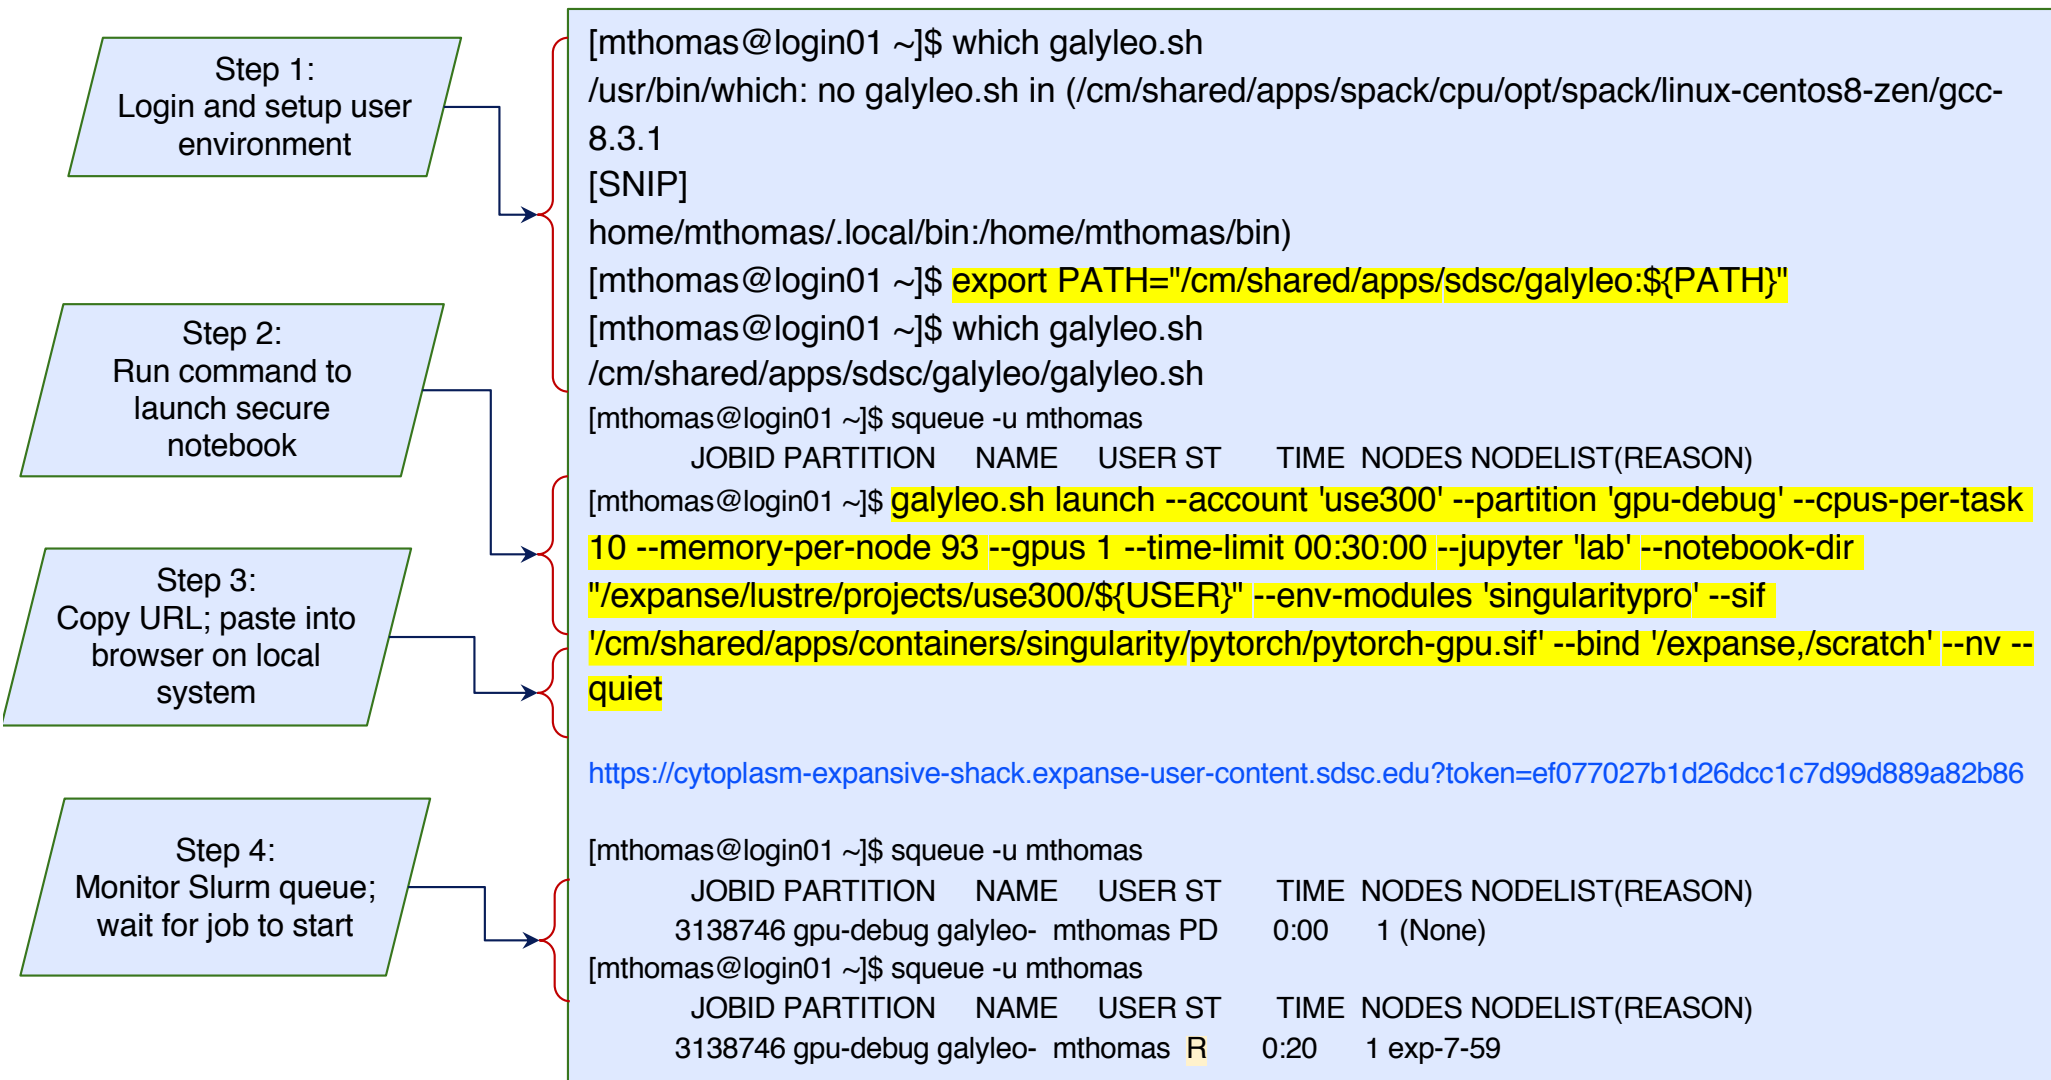
```
# Example 1: Launch a Jupyter Notebook session on a single CPU core in the 'debug'  
# partition on Expanse using the 'base' Anaconda3 software environment provided as part  
# of Expanse's standard software modules.
```

```
/cm/shared/apps/sdsc/galileo/galileo.sh launch --account 'use300' --partition 'debug' --  
cpus-per-task 1 --memory-per-node 1 --time-limit 00:30:00 --jupyter 'notebook' --notebook-  
dir "/exppanse/lustre/projects/use300/${USER}" --env-modules 'cpu,gcc,anaconda3' --  
conda-env 'base' --quiet
```

```
# Example 2: Launch a JupyterLab session on a single GPU in the 'gpu-debug' partition  
# on Expanse using the latest PyTorch Singularity container available.
```

```
galileo.sh launch --account 'use300' --partition 'gpu-debug' --cpus-per-task 10 --memory-per-node 93 --  
gpus 1 --time-limit 00:30:00 --jupyter 'lab' --notebook-dir "/exppanse/lustre/projects/use300/${USER}" --  
env-modules 'singularitypro' --sif '/cm/shared/apps/containers/singularity/pytorch/pytorch-gpu.sif' --bind  
'/exppanse,/scratch' --nv --quiet
```

Running GPU notebooks using galyleo.sh



- Paste the **HTTPS** URL into a web browser

Satellite Service: Monitor Notebook Status

- Load notebook URL in browser; wait for it to launch
- Monitor pending page
- Run the “squeue” command on the HPC system to check job status
- If the job queue is busy, it may take a while to launch the notebook
- Treat Jupyter Notebook URL as a password

Satellite Reverse Proxy Service

SDSC Expanse

Job State: Mapped



In Queue

Job has not yet started.

Running

Job has started, but has not redeemed Satellite Token.

Mapped

Job has redeemed Satellite Token, but no proxy entry exists yet.

Proxied

Proxy entry created, ready to go!

Dead

Job died or exited, no further progress will occur.

Satellite Pending Page

Satellite Reverse Proxy Service

SDSC Expanse

Copy the URL and paste into Web browser

Job State: Unknown

In Queue Running Mapped Proxied

Satellite Reverse Proxy Service

SDSC Expanse

Job State: Mapped

In Queue Running Mapped Proxied

Monitor status window: may take some time

Satellite Reverse Proxy Service

SDSC Expanse

Job State: Proxied

In Queue Running Mapped Proxied

In Queue
Job has not yet started.

Running
Job has started, but has not redeemed Satellite Token.

Mapped
Job has redeemed Satellite Token, but no proxy entry exists yet.

Proxied
Proxy entry created, ready to go!

Dead
Job died or exited, no further progress will occur.

Access your **secure** Jupyter Service when compute node is allocated

dipping-either-pureblood.expance-user-content.sdsc.edu/lab

File Edit View Run Kernel Tabs Settings Help

numpy_intro.ipynb

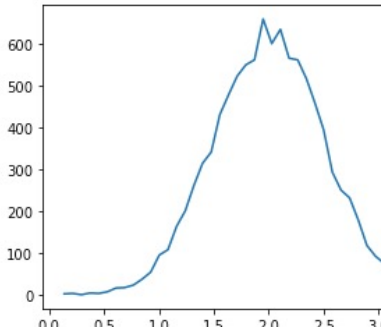
examples / NumPy_Intro /

Name

numpy_intro.ipynb

README.md

```
plt.plot(.5*(bin_edges[1:]+bin_edges[2:]),  
plt.show())
```



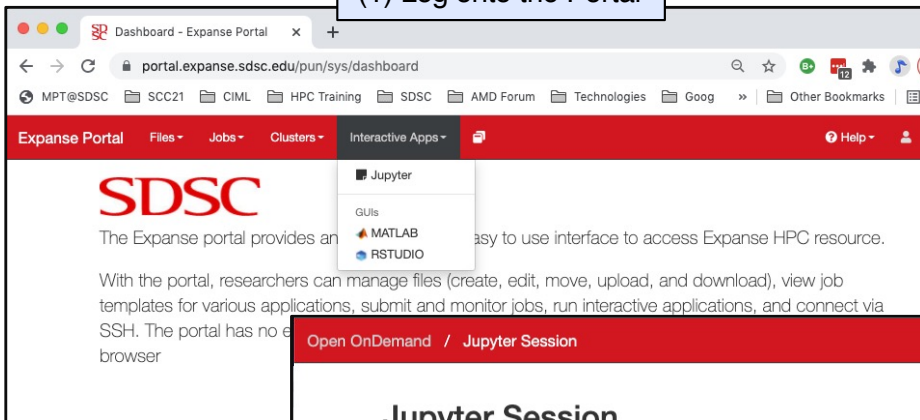
Linear algebra

NumPy provides a comprehensive set of linear algebra functions. These are listed below. A convenient summary of these functions is available at <https://docs.scipy.org/doc/numpy-1.13.0/reference/arrays.linalg.html>.

In this first example, we calculate the inverse of a matrix.

Running Notebooks from Expanse Portal

(1) Log onto the Portal



(2) Fill out form inputs

Jupyter Session

Account:
sds184

Partition (Please choose the gpu, gpu-shared, or gpu-preempt as the partition if using gpus):
shared

Time limit (min):
30

Number of cores:
1

Memory required per node (GB):
2

GPUs (optional):

(3) Copy the URL and paste into Web browser



Satellite Reverse Proxy Service

SDSC Expanse

(4) Monitor status window:
may take a long time

Job State: Proxied



In Queue
Job has not yet started.

Running
Job has started, but has not redeemed Satellite Token.

A screenshot of a Jupyter Notebook interface. The left sidebar shows a file explorer with 'numpy_intro.ipynb' selected. The main area displays a plot of a bell curve (Gaussian distribution) with the x-axis ranging from 0.0 to 4.0 and the y-axis from 0 to 600. Below the plot, there is a section titled 'Linear algebra' with text explaining NumPy's capabilities and a code snippet for matrix inversion. The code includes comments and uses NumPy functions like 'np.random.rand', 'np.linalg.inv', and 'np.eye'.

(5) Access your
Jupyter Service

Satellite Client

Example:

`start-jupyter`

Running start-notebook

Make sure you have activated your conda environment

```
(base) [mthomas@login02 reverse-proxy]$ ./start-jupyter -A use300
Your notebook is here:
https://tightness-pulverize-starting.expanse-user-content.sdsc.edu?token=68a79db47927f40bc3831ce0c37f87a3
If you encounter any issues, please email help@xsede.org and mention the Reverse Proxy Service.
If you're not on Xsede, please email consult@sdsc.edu instead.
Your job id is 4543288
You may occasionally run the command 'squeue -j 4543288' to check the status of your job
(base) [mthomas@login02 reverse-proxy]$ squeue -j 4543288
      JOBID PARTITION  NAME  USER ST  TIME  NODES NODELIST(REASON)
      4543288  compute notebook  mthomas  R   0:50    1 exp-6-32
```

Once notebook is running, you can check the slurm output file

```
(base) [mthomas@login02 reverse-proxy]$ cat slurm-4543288.out
Image:
[I 01:52:43.924 NotebookApp] JupyterLab extension loaded from /cm/shared/apps/spack/gpu/opt/spack/linux-centos8-skylake_avx512/gcc-8.3.1/anaconda3-2020.11-bsn4npoxyw7jzz7fajncek3bvdoaa5wv/lib/python3.8/site-packages/jupyterlab
[I 01:52:43.925 NotebookApp] JupyterLab application directory is /cm/shared/apps/spack/gpu/opt/spack/linux-centos8-skylake_avx512/gcc-8.3.1/anaconda3-2020.11-bsn4npoxyw7jzz7fajncek3bvdoaa5wv/share/jupyter/lab
[I 01:52:43.926 NotebookApp] Serving notebooks from local directory: /home/mthomas
[I 01:52:43.926 NotebookApp] Jupyter Notebook 6.1.4 is running at:
[I 01:52:43.926 NotebookApp] http://exp-6-32.eth.cluster:8888/?token=...
[I 01:52:43.926 NotebookApp] or http://127.0.0.1:8888/?token=...
[I 01:52:43.926 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).

 % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent  Left  Speed
100    9 100    9  0    0   53    0 --:--:-- --:--:-- --:--:--   52
Success!
[I 01:53:21.097 NotebookApp] 302 GET /?token=68a79db47927f40bc3831ce0c37f87a3 (10.21.0.30) 0.53ms
```

Paste the HTTPS URL into a web browser

Satellite Server Pending Page

- Load notebook URL in browser; wait for it to launch
- Monitor pending page
- Run the “squeue” command on the HPC system to check job status
- If the job queue is busy, it may take a while to launch the notebook
- Treat Jupyter Notebook URL as a password

Satellite Reverse Proxy Service

SDSC Expanse



Job State: Mapped



In Queue

Job has not yet started.

Running

Job has started, but has not redeemed Satellite Token.

Mapped

Job has redeemed Satellite Token, but no proxy entry exists yet.

Proxied

Proxy entry created, ready to go!

Dead

Job died or exited, no further progress will occur.

Your notebook is launched

The screenshot shows a web browser window with the address bar containing the URL `annuity-headphone-aptitude.expense-user-content.sdsc.edu/tree`. The browser's address bar and the JupyterLab interface's 'Quit' button are circled in yellow. The JupyterLab interface displays the 'Files' tab, showing a list of files and folders in a table format. The table has columns for 'Name', 'Last Modified', and 'File size'.

Name	Last Modified	File size
conda-install-tmp	a month ago	
DEMO	11 hours ago	
dev	2 months ago	
hpctrain	2 months ago	
matlab.dev	a month ago	
miniconda3	6 days ago	
notebook-examples	6 days ago	
rev-pxy	12 minutes ago	
tools	2 months ago	
conda-activate.txt	29 minutes ago	224 B
demo.tar	2 months ago	3.32 MB
modules.hpc.sh	a month ago	61 B
README	2 months ago	124 B

When done with the notebook be sure to **shut it down** by quitting the notebook

Thank You

Resources

- Training Material for this presentation:
 - <https://hpc-training.sdsc.edu/>
 - <https://github.com/sdsc-hpc-training-org/expanse-101>
 - <https://github.com/sdsc-hpc-training-org/notebooks-101>
 - Interactive tutorial:
https://education.sdsc.edu/training/interactive/202012_running_jupyter_notebooks_on_expanse/index.html
- Expanse User Guide
 - https://www.sdsc.edu/support/user_guides/expanse.html
- SDSC Training Resources
 - https://www.sdsc.edu/education_and_training/training
 - <https://github.com/sdsc-hpc-training/webinars>
- XSEDE Training Resources
 - <https://www.xsede.org/for-users/training>
 - <https://cvw.cac.cornell.edu/expanse/>